

# CPSC 540: Machine Learning

## Kernel Density Estimation, Factor Analysis

Mark Schmidt

University of British Columbia

Winter 2017

# Admin

- **Assignment 2:**
  - 2 late days to hand it in tonight.
- **Assignment 3:**
  - Due February 27.

## Comments on TensorFlow Talk

- Most of the talk focused on large-scale issues, which I won't cover:
  - Synchronous vs. asynchronous (540 course project topic in 2014).
  - Issues related to distributed data/parameters.
- Some models were mentioned that I'm planning to get to:
  - Word2vec.
  - RNNs.
  - LSTMs.
  - Sequence-to-sequence.
  - Neural machine translation.

## Last Time: Mixture of Gaussians

- The classic **mixture of Gaussians** model uses a PDF of the form

$$p(x^i|\Theta) = \sum_{c=1}^k \underbrace{p(z^i = c|\Theta)}_{\text{prob(cluster)}} \underbrace{p(x^i|z^i = c, \Theta)}_{\text{prob}(x) \text{ given cluster}},$$

## Last Time: Mixture of Gaussians

- The classic **mixture of Gaussians** model uses a PDF of the form

$$p(x^i|\Theta) = \sum_{c=1}^k \underbrace{p(z^i = c|\Theta)}_{\text{prob(cluster)}} \underbrace{p(x^i|z^i = c, \Theta)}_{\text{prob}(x) \text{ given cluster}},$$

where each mixture component is a multivariate Gaussian,

$$p(x^i|z^i = c, \Theta) = \frac{1}{(2\pi)^{\frac{d}{2}} |\Sigma_c|^{\frac{1}{2}}} \exp \left( -\frac{1}{2} (x^i - \mu_c)^T \Sigma_c^{-1} (x^i - \mu_c) \right),$$

## Last Time: Mixture of Gaussians

- The classic **mixture of Gaussians** model uses a PDF of the form

$$p(x^i|\Theta) = \sum_{c=1}^k \underbrace{p(z^i = c|\Theta)}_{\text{prob(cluster)}} \underbrace{p(x^i|z^i = c, \Theta)}_{\text{prob}(x) \text{ given cluster}},$$

where each mixture component is a multivariate Gaussian,

$$p(x^i|z^i = c, \Theta) = \frac{1}{(2\pi)^{\frac{d}{2}} |\Sigma_c|^{\frac{1}{2}}} \exp \left( -\frac{1}{2} (x^i - \mu_c)^T \Sigma_c^{-1} (x^i - \mu_c) \right),$$

and we model the mixture probabilities as categorical,

$$p(z^i = c|\Theta) = \pi_c.$$

## Last Time: Mixture of Gaussians

- The classic **mixture of Gaussians** model uses a PDF of the form

$$p(x^i|\Theta) = \sum_{c=1}^k \underbrace{p(z^i = c|\Theta)}_{\text{prob(cluster)}} \underbrace{p(x^i|z^i = c, \Theta)}_{\text{prob}(x) \text{ given cluster}},$$

where each mixture component is a multivariate Gaussian,

$$p(x^i|z^i = c, \Theta) = \frac{1}{(2\pi)^{\frac{d}{2}} |\Sigma_c|^{\frac{1}{2}}} \exp \left( -\frac{1}{2} (x^i - \mu_c)^T \Sigma_c^{-1} (x^i - \mu_c) \right),$$

and we model the mixture probabilities as categorical,

$$p(z^i = c|\Theta) = \pi_c.$$

- Finding the optimal parameter  $\Theta = \{\pi_c, \mu_c, \Sigma_c\}_{c=1}^k$  is NP-hard.
  - But EM updates for improving parameters use analytic form of Gaussian MLE.

## Expectation Maximization for Mixture of Gaussians

- EM update for mixture models is often written in terms of **responsibilities**,

$$r_c^i \triangleq p(z^i = c | x^i, \Theta^t) = \frac{p(x^i | z^i = c, \Theta^t) p(z^i = c, \Theta^t)}{\sum_{c'=1}^k p(x^i | z^i = c', \Theta^t) p(z^i = c', \Theta^t)},$$

a number between 0 and 1 giving probability that cluster  $c$  generated  $x^i$ .



## Expectation Maximization for Mixture of Gaussians

- EM update for mixture models is often written in terms of **responsibilities**,

$$r_c^i \triangleq p(z^i = c | x^i, \Theta^t) = \frac{p(x^i | z^i = c, \Theta^t) p(z^i = c, \Theta^t)}{\sum_{c'=1}^k p(x^i | z^i = c', \Theta^t) p(z^i = c', \Theta^t)},$$

a number between 0 and 1 giving probability that cluster  $c$  generated  $x^i$ .

- For mixture of Gaussian, EM updates takes the form

$$\pi_c^{t+1} = \frac{1}{n} \sum_{i=1}^n r_c^i \quad \text{(proportion of examples soft-assigned to cluster } c)$$

$$\mu_c^{t+1} = \frac{\sum_{i=1}^n r_c^i x^i}{n \pi_c^{t+1}} \quad \text{(mean of examples soft-assigned to cluster } c)$$

$$\Sigma_c^{t+1} = \frac{\sum_{i=1}^n r_c^i (x^i - \mu_c^{t+1})(x^i - \mu_c^{t+1})^T}{n \pi_c^{t+1}} \quad \text{(covariance of examples soft-assigned to } c).$$

## Expectation Maximization for Mixture of Gaussians

- EM update for mixture models is often written in terms of **responsibilities**,

$$r_c^i \triangleq p(z^i = c | x^i, \Theta^t) = \frac{p(x^i | z^i = c, \Theta^t) p(z^i = c, \Theta^t)}{\sum_{c'=1}^k p(x^i | z^i = c', \Theta^t) p(z^i = c', \Theta^t)},$$

a number between 0 and 1 giving probability that cluster  $c$  generated  $x^i$ .

- For mixture of Gaussian, EM updates takes the form

$$\pi_c^{t+1} = \frac{1}{n} \sum_{i=1}^n r_c^i \quad (\text{proportion of examples soft-assigned to cluster } c)$$

$$\mu_c^{t+1} = \frac{\sum_{i=1}^n r_c^i x^i}{n \pi_c^{t+1}} \quad (\text{mean of examples soft-assigned to cluster } c)$$

$$\Sigma_c^{t+1} = \frac{\sum_{i=1}^n r_c^i (x^i - \mu_c^{t+1})(x^i - \mu_c^{t+1})^T}{n \pi_c^{t+1}} \quad (\text{covariance of examples soft-assigned to } c).$$

- Derivation is tedious (see notes on webpage).
  - Uses distributive law, probabilities sum to one, Lagrangian, weighted Gaussian MLE.
- We get  $k$ -means if  $r_c^i = 1$  for most likely cluster, and  $\Sigma_c$  is constant across  $c$ .

# Expectation Maximization for Mixture of Gaussians

- EM for fitting mixture of Gaussians in action:  
<https://www.youtube.com/watch?v=B36fzChfyGU>


## K-Means vs. Mixture of Gaussians

- K-means is special case of “hard EM” for mixture of Gaussians (common  $\Sigma_c$ ).
  - But EM allows points to be assigned to be multiple clusters

## K-Means vs. Mixture of Gaussians


- K-means is special case of “hard EM” for mixture of Gaussians (common  $\Sigma_c$ ).
  - But EM allows points to be assigned to be multiple clusters
  - General  $\Sigma_c$  in mixture of Gaussians allow **non-convex clusters**.

With same covariance, clusters are convex.




## K-Means vs. Mixture of Gaussians

- K-means is special case of “hard EM” for mixture of Gaussians (common  $\Sigma_c$ ).
  - But EM allows points to be assigned to be multiple clusters
  - General  $\Sigma_c$  in mixture of Gaussians allow **non-convex clusters**.




## K-Means vs. Mixture of Gaussians

- K-means is special case of “hard EM” for mixture of Gaussians (common  $\Sigma_c$ ).
  - But EM allows points to be assigned to be multiple clusters
  - General  $\Sigma_c$  in mixture of Gaussians allow **non-convex clusters**.



## K-Means vs. Mixture of Gaussians

- K-means is special case of “hard EM” for mixture of Gaussians (common  $\Sigma_c$ ).
  - But EM allows points to be assigned to be multiple clusters
  - General  $\Sigma_c$  in mixture of Gaussians allow **non-convex clusters**.





# Outline

- 1 Monotonicity of EM
- 2 Kernel Density Estimation
- 3 Factor Analysis

# Expectation Maximization

- EM considers learning with **observed variables**  $O$  and **hidden variables**  $H$ .
- In this case the “observed” marginal log-likelihood has a nasty form,

$$\log p(O|\Theta) = \log \left( \sum_H p(O, H|\Theta) \right).$$

## Expectation Maximization

- EM considers learning with **observed variables**  $O$  and **hidden variables**  $H$ .
- In this case the “observed” marginal log-likelihood has a nasty form,

$$\log p(O|\Theta) = \log \left( \sum_H p(O, H|\Theta) \right).$$

- **EM** applies when “complete” likelihood,  $p(O, H|\Theta)$ , has a nice form.
- EM iterations take the form

$$\Theta^{t+1} = \operatorname{argmax}_{\Theta} \left\{ \sum_H \alpha_H \log p(O, H|\Theta) \right\},$$

where  $\alpha_H = p(H|O, \Theta^t)$ .

## Bound on Progress of Expectation Maximization

*The iterations of the EM algorithm satisfy*

$$\log p(O|\Theta^{t+1}) - \log p(O|\Theta^t) \geq Q(\Theta^{t+1}|\Theta^t) - Q(\Theta^t|\Theta^t),$$

## Bound on Progress of Expectation Maximization

*The iterations of the EM algorithm satisfy*

$$\log p(O|\Theta^{t+1}) - \log p(O|\Theta^t) \geq Q(\Theta^{t+1}|\Theta^t) - Q(\Theta^t|\Theta^t),$$

• Proof:

$$\begin{aligned} -\log p(O|\Theta) &= -\log \left( \sum_h p(O, H|\Theta) \right) \\ &= -\log \left( \sum_H \alpha_H \frac{p(O, H|\Theta)}{\alpha_H} \right) \quad (\text{for } \alpha_H \neq 0) \end{aligned}$$

## Bound on Progress of Expectation Maximization

*The iterations of the EM algorithm satisfy*

$$\log p(O|\Theta^{t+1}) - \log p(O|\Theta^t) \geq Q(\Theta^{t+1}|\Theta^t) - Q(\Theta^t|\Theta^t),$$

• Proof:

$$\begin{aligned} -\log p(O|\Theta) &= -\log \left( \sum_h p(O, H|\Theta) \right) \\ &= -\log \left( \sum_H \alpha_H \frac{p(O, H|\Theta)}{\alpha_H} \right) && \text{(for } \alpha_H \neq 0) \\ &\leq -\sum_H \alpha_H \log \left( \frac{p(O, H|\Theta)}{\alpha_H} \right), \end{aligned}$$

because  $-\log(z)$  is convex.

## Bound on Progress of Expectation Maximization

*The iterations of the EM algorithm satisfy*

$$\log p(O|\Theta^{t+1}) - \log p(O|\Theta^t) \geq Q(\Theta^{t+1}|\Theta^t) - Q(\Theta^t|\Theta^t),$$

- Using that log turns multiplication into addition we get

$$\begin{aligned} -\log p(O|\Theta) &\leq -\sum_H \alpha_H \log \left( \frac{p(O, H|\Theta)}{\alpha_H} \right) \\ &= -\underbrace{\sum_H \alpha_H \log p(O, H|\Theta)}_{Q(\Theta|\Theta^t)} + \underbrace{\sum_H \alpha_H \log \alpha_H}_{\text{negative entropy}} \end{aligned}$$

## Bound on Progress of Expectation Maximization

*The iterations of the EM algorithm satisfy*

$$\log p(O|\Theta^{t+1}) - \log p(O|\Theta^t) \geq Q(\Theta^{t+1}|\Theta^t) - Q(\Theta^t|\Theta^t),$$

- Using that log turns multiplication into addition we get

$$\begin{aligned} -\log p(O|\Theta) &\leq -\sum_H \alpha_H \log \left( \frac{p(O, H|\Theta)}{\alpha_H} \right) \\ &= -\underbrace{\sum_H \alpha_H \log p(O, H|\Theta)}_{Q(\Theta|\Theta^t)} + \underbrace{\sum_H \alpha_H \log \alpha_H}_{\text{negative entropy}} \\ &= -Q(\Theta|\Theta^t) - \text{entropy}(\alpha), \end{aligned}$$

which we can use to bound  $\log p(O|\Theta^{t+1})$ .



## Bound on Progress of Expectation Maximization

*The iterations of the EM algorithm satisfy*

$$\log p(O|\Theta^{t+1}) - \log p(O|\Theta^t) \geq Q(\Theta^{t+1}|\Theta^t) - Q(\Theta^t|\Theta^t),$$

- To bound  $p(O|\Theta^t)$  we use definition of conditional,

$$p(H|O, \Theta^t) = \frac{p(O, H|\Theta^t)}{p(O|\Theta^t)}$$

## Bound on Progress of Expectation Maximization

*The iterations of the EM algorithm satisfy*

$$\log p(O|\Theta^{t+1}) - \log p(O|\Theta^t) \geq Q(\Theta^{t+1}|\Theta^t) - Q(\Theta^t|\Theta^t),$$

- To bound  $p(O|\Theta^t)$  we use definition of conditional,

$$p(H|O, \Theta^t) = \frac{p(O, H|\Theta^t)}{p(O|\Theta^t)} \quad \text{or} \quad \log p(O|\Theta^t) = \log p(O, H|\Theta^t) - \log p(H|O, \Theta^t).$$

## Bound on Progress of Expectation Maximization

*The iterations of the EM algorithm satisfy*

$$\log p(O|\Theta^{t+1}) - \log p(O|\Theta^t) \geq Q(\Theta^{t+1}|\Theta^t) - Q(\Theta^t|\Theta^t),$$

- To bound  $p(O|\Theta^t)$  we use definition of conditional,

$$p(H|O, \Theta^t) = \frac{p(O, H|\Theta^t)}{p(O|\Theta^t)} \quad \text{or} \quad \log p(O|\Theta^t) = \log p(O, H|\Theta^t) - \log p(H|O, \Theta^t).$$

- Multiply by  $\alpha_H$  and sum equality over  $H$  values,

$$\sum_H \alpha_H \log p(O|\Theta^t) = \sum_H \alpha_H \log p(O, H|\Theta^t) - \sum_H \alpha_H \log p(H|O, \Theta^t).$$

## Bound on Progress of Expectation Maximization

*The iterations of the EM algorithm satisfy*

$$\log p(O|\Theta^{t+1}) - \log p(O|\Theta^t) \geq Q(\Theta^{t+1}|\Theta^t) - Q(\Theta^t|\Theta^t),$$

- To bound  $p(O|\Theta^t)$  we use definition of conditional,

$$p(H|O, \Theta^t) = \frac{p(O, H|\Theta^t)}{p(O|\Theta^t)} \quad \text{or} \quad \log p(O|\Theta^t) = \log p(O, H|\Theta^t) - \log p(H|O, \Theta^t).$$

- Multiply by  $\alpha_H$  and sum equality over  $H$  values,

$$\sum_H \alpha_H \log p(O|\Theta^t) = \sum_H \alpha_H \log p(O, H|\Theta^t) - \sum_H \alpha_H \log p(H|O, \Theta^t).$$

- Using the EM definition of  $\alpha_h$  we have

$$\log p(O|\Theta^t) \underbrace{\sum_H \alpha_H}_{=1} = Q(\Theta^t|\Theta^t) + \text{entropy}(\alpha).$$

## Bound on Progress of Expectation Maximization

*The iterations of the EM algorithm satisfy*

$$\log p(O|\Theta^{t+1}) - \log p(O|\Theta^t) \geq Q(\Theta^{t+1}|\Theta^t) - Q(\Theta^t|\Theta^t),$$

- Thus we have the two bounds

$$\log p(O|\Theta) \geq Q(\Theta|\Theta^t) + \text{entropy}(\alpha)$$

$$\log p(O|\Theta^t) = Q(\Theta^t|\Theta^t) + \text{entropy}(\alpha).$$

Subtracting these and using  $\Theta = \Theta^{t+1}$  gives the result.

## Bound on Progress of Expectation Maximization

*The iterations of the EM algorithm satisfy*

$$\log p(O|\Theta^{t+1}) - \log p(O|\Theta^t) \geq Q(\Theta^{t+1}|\Theta^t) - Q(\Theta^t|\Theta^t),$$

- Thus we have the two bounds

$$\log p(O|\Theta) \geq Q(\Theta|\Theta^t) + \text{entropy}(\alpha)$$

$$\log p(O|\Theta^t) = Q(\Theta^t|\Theta^t) + \text{entropy}(\alpha).$$

Subtracting these and using  $\Theta = \Theta^{t+1}$  gives the result.

- Inequality holds for any choice of  $\Theta^{t+1}$ .
  - **Approximate M-steps are ok:** we just need to decrease  $Q$  to improve likelihood.
- Implies **entropy of  $\alpha_H$**  gives tightness of bound.
  - If variables are “predictable” then the bound is tight and we get “hard” EM.

## Discussing of EM for Mixtures of Gaussians

- EM and mixture models are used in a ton of applications.
  - One of the default unsupervised learning methods.

## Discussing of EM for Mixtures of Gaussians

- EM and mixture models are used in a ton of applications.
  - One of the default unsupervised learning methods.
- EM usually doesn't reach global optimum.
  - Classic solution: restart the algorithm from different initializations.



## Discussing of EM for Mixtures of Gaussians

- EM and mixture models are used in a ton of applications.
  - One of the default unsupervised learning methods.
- EM usually doesn't reach global optimum.
  - Classic solution: restart the algorithm from different initializations.
- MLE for some clusters may not exist (e.g., only responsible for one point).
  - Use MAP estimates or remove these clusters.

## Discussing of EM for Mixtures of Gaussians

- EM and mixture models are used in a ton of applications.
  - One of the default unsupervised learning methods.
- EM usually doesn't reach global optimum.
  - Classic solution: restart the algorithm from different initializations.
- MLE for some clusters may not exist (e.g., only responsible for one point).
  - Use MAP estimates or remove these clusters.
- How do you choose number of mixtures  $k$ ?
  - Use cross-validation or other model selection criteria.

## Discussing of EM for Mixtures of Gaussians

- EM and mixture models are used in a ton of applications.
  - One of the default unsupervised learning methods.
- EM usually doesn't reach global optimum.
  - Classic solution: restart the algorithm from different initializations.
- MLE for some clusters may not exist (e.g., only responsible for one point).
  - Use MAP estimates or remove these clusters.
- How do you choose number of mixtures  $k$ ?
  - Use cross-validation or other model selection criteria.
- Can you make it robust?
  - Use mixture of Laplace or student  $t$  distributions.

## Discussing of EM for Mixtures of Gaussians

- EM and mixture models are used in a ton of applications.
  - One of the default unsupervised learning methods.
- EM usually doesn't reach global optimum.
  - Classic solution: restart the algorithm from different initializations.
- MLE for some clusters may not exist (e.g., only responsible for one point).
  - Use MAP estimates or remove these clusters.
- How do you choose number of mixtures  $k$ ?
  - Use cross-validation or other model selection criteria.
- Can you make it robust?
  - Use mixture of Laplace or student  $t$  distributions.
- Are there alternatives to EM?
  - Could use gradient descent on NLL.
  - [Spectral](#) and other recent methods have some global guarantees.

# Outline

- 1 Monotonicity of EM
- 2 Kernel Density Estimation
- 3 Factor Analysis

## A Non-Parametric Mixture Model

- The classic **parametric** mixture model has the form

$$p(x) = \sum_{c=1}^k p(z = c)p(x|z = c).$$

## A Non-Parametric Mixture Model

- The classic **parametric** mixture model has the form

$$p(x) = \sum_{c=1}^k p(z = c)p(x|z = c).$$

- A natural way to define a **non-parametric** mixture model is

$$p(x) = \sum_{i=1}^n p(z = i)p(x|z = i),$$

where we have **one mixture for every training example  $i$** .

## A Non-Parametric Mixture Model

- The classic **parametric** mixture model has the form

$$p(x) = \sum_{c=1}^k p(z = c)p(x|z = c).$$

- A natural way to define a **non-parametric** mixture model is

$$p(x) = \sum_{i=1}^n p(z = i)p(x|z = i),$$

where we have **one mixture for every training example  $i$** .

- Common example:  $z$  is uniform and  $x|z$  is Gaussian with mean  $x^i$ ,

$$p(x) = \frac{1}{n} \sum_{i=1}^n \mathcal{N}(x|x^i, \sigma^2 I),$$

and we use a shared covariance  $\sigma^2 I$  (with  $\sigma$  estimated by cross-validation).



## A Non-Parametric Mixture Model

- The classic **parametric** mixture model has the form

$$p(x) = \sum_{c=1}^k p(z = c)p(x|z = c).$$

- A natural way to define a **non-parametric** mixture model is

$$p(x) = \sum_{i=1}^n p(z = i)p(x|z = i),$$

where we have **one mixture for every training example  $i$** .

- Common example:  $z$  is uniform and  $x|z$  is Gaussian with mean  $x^i$ ,


$$p(x) = \frac{1}{n} \sum_{i=1}^n \mathcal{N}(x|x^i, \sigma^2 I),$$

and we use a shared covariance  $\sigma^2 I$  (with  $\sigma$  estimated by cross-validation).

- This is a special case of **kernel density estimation** or **Parzen window**.



# Histogram vs. Kernel Density Estimator

- Think of kernel density estimator as a smooth version of **histogram**:







[https://en.wikipedia.org/wiki/Kernel\\_density\\_estimation](https://en.wikipedia.org/wiki/Kernel_density_estimation)





# Parzen Window vs. Gaussian and Mixture of Gaussian



# Parzen Window vs. Gaussian and Mixture of Gaussian



# Parzen Window vs. Gaussian and Mixture of Gaussian



# Kernel Density Estimation

- The 1D kernel density estimation (KDE) model uses

$$p(x) = \frac{1}{n} \sum_{i=1}^n k_{\sigma}(x - x^i),$$

where the PDF  $k$  is the “kernel” and the parameter  $\sigma$  is the “bandwidth”.

## Kernel Density Estimation

- The 1D kernel density estimation (KDE) model uses

$$p(x) = \frac{1}{n} \sum_{i=1}^n k_{\sigma}(x - x^i),$$

where the PDF  $k$  is the “kernel” and the parameter  $\sigma$  is the “bandwidth”.

- In the previous slide we used the (normalized) Gaussian kernel,

$$k_1(x) = \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{x^2}{2}\right), \quad k_{\sigma}(x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{x^2}{2\sigma^2}\right).$$

# Kernel Density Estimation

- The 1D **kernel density estimation** (KDE) model uses

$$p(x) = \frac{1}{n} \sum_{i=1}^n k_{\sigma}(x - x^i),$$

where the PDF  $k$  is the “**kernel**” and the parameter  $\sigma$  is the “**bandwidth**”.

- In the previous slide we used the (normalized) Gaussian kernel,

$$k_1(x) = \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{x^2}{2}\right), \quad k_{\sigma}(x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{x^2}{2\sigma^2}\right).$$

- Note that we can add a bandwidth  $\sigma$  to any PDF  $k_1$ , using

$$k_{\sigma}(x) = \frac{1}{\sigma} k_1\left(\frac{x}{\sigma}\right),$$

which follows from the **change of variables** formula for probabilities.

- Under common choices of kernels, **KDEs** can model any continuous density.



# Efficient Kernel Density Estimation

- KDE with the Gaussian kernel is **slow at test time**:
  - We need to compute distance of test point to every training point.

## Efficient Kernel Density Estimation

- KDE with the Gaussian kernel is **slow at test time**:
  - We need to compute distance of test point to every training point.
- A common alternative is the **Epanechnikov** kernel,

$$k_1(x) = \frac{3}{4} (1 - x^2) \mathcal{I}[|x| \leq 1].$$

- This kernel has two nice properties:
  - Epanechnikov showed that it is **asymptotically optimal** in terms of squared error.
  - It can be much faster to use since it only depends on nearby points.
    - You can use fast methods for computing nearest neighbours.

## Efficient Kernel Density Estimation



- KDE with the Gaussian kernel is **slow at test time**:
  - We need to compute distance of test point to every training point.
- A common alternative is the **Epanechnikov** kernel,

$$k_1(x) = \frac{3}{4} (1 - x^2) \mathcal{I}[|x| \leq 1] .$$

- This kernel has two nice properties:
  - Epanechnikov showed that it is **asymptotically optimal** in terms of squared error.
  - It can be much faster to use since it only depends on nearby points.
    - You can use fast methods for computing nearest neighbours.
- It is non-smooth at the boundaries but many smooth approximations exist.
  - Quartic, triweight, tricube, cosine, etc.




# Visualization of Common Kernel Functions

Histogram vs. Gaussian vs. Epanechnikov vs. tricube:







# Visualization of Common Kernel Functions

Histogram vs. Gaussian vs. Epanechnikov vs. tricube:



# Visualization of Common Kernel Functions

Histogram vs. Gaussian vs. Epanechnikov vs. tricube:



## Multivariate Kernel Density Estimation

- The multivariate **kernel density estimation** (KDE) model uses

$$p(x) = \frac{1}{n} \sum_{i=1}^n k_R(x - x^i),$$

- The most common kernel is again the Gaussian,

$$k_I(x) = \frac{1}{\sqrt{2\pi}^{\frac{d}{2}}} \exp\left(-\frac{\|x\|^2}{2}\right).$$

## Multivariate Kernel Density Estimation

- The multivariate **kernel density estimation** (KDE) model uses

$$p(x) = \frac{1}{n} \sum_{i=1}^n k_R(x - x^i),$$

- The most common kernel is again the Gaussian,

$$k_I(x) = \frac{1}{\sqrt{2\pi}^{\frac{d}{2}}} \exp\left(-\frac{\|x\|^2}{2}\right).$$

- We can add a **bandwith matrix**  $R$  to any kernel using

$$k_R(x) = \frac{1}{|R|} k_1(R^{-1}x) \quad \left(\text{generalizes } k_\sigma(x) = \frac{1}{\sigma} k_1\left(\frac{x}{\sigma}\right)\right),$$

and multivariate Gaussian with covariance  $\Sigma$  corresponds to  $R = \Sigma^{\frac{1}{2}}$ .

- To reduce number of parameters, we typically:
  - Use a product of independent distributions and use  $R = \sigma I$  for some  $\sigma$ .



# Mean-Shift Clustering

- Mean-shift clustering uses KDE for clustering:
  - Fit a KDE to the training examples, and then for test example  $\hat{x}$ :
    - Run gradient descent starting from  $\hat{x}$ .
  - Clusters are points that reach same local minimum.
- <https://spin.atomicobject.com/2015/05/26/mean-shift-clustering>

# Mean-Shift Clustering

- Mean-shift clustering uses KDE for clustering:
  - Fit a KDE to the training examples, and then for test example  $\hat{x}$ :
    - Run gradient descent starting from  $\hat{x}$ .
    - Clusters are points that reach same local minimum.
- <https://spin.atomicobject.com/2015/05/26/mean-shift-clustering>
- Not sensitive to initialization, no need to choose  $k$ , can find non-convex clusters.
- Similar to density-based clustering from 340.
  - But doesn't require uniform density within cluster.
  - And can be used for vector quantization.

# Outline

- 1 Monotonicity of EM
- 2 Kernel Density Estimation
- 3 Factor Analysis

## Expectation Maximization with Many Discrete Variables

- EM iterations take the form

$$\Theta^{t+1} = \operatorname{argmax}_{\Theta} \left\{ \sum_H \alpha_H \log p(O, H | \Theta) \right\},$$

and with multiple MAR variables  $\{H_1, H_2, \dots, H_m\}$  this means

$$\Theta^{t+1} = \operatorname{argmax}_{\Theta} \left\{ \sum_{H_1} \sum_{H_2} \cdots \sum_{H_m} \alpha_H \log p(O, H | \Theta) \right\},$$

## Expectation Maximization with Many Discrete Variables

- EM iterations take the form

$$\Theta^{t+1} = \operatorname{argmax}_{\Theta} \left\{ \sum_H \alpha_H \log p(O, H | \Theta) \right\},$$

and with multiple MAR variables  $\{H_1, H_2, \dots, H_m\}$  this means

$$\Theta^{t+1} = \operatorname{argmax}_{\Theta} \left\{ \sum_{H_1} \sum_{H_2} \cdots \sum_{H_m} \alpha_H \log p(O, H | \Theta) \right\},$$

- In mixture models, EM **sums over all  $2^n$**  possible cluster assignments.
- In binary semi-supervised learning, EM **sums over all  $2^t$**  assignments to  $\tilde{y}$ .

## Expectation Maximization with Many Discrete Variables

- EM iterations take the form

$$\Theta^{t+1} = \operatorname{argmax}_{\Theta} \left\{ \sum_H \alpha_H \log p(O, H | \Theta) \right\},$$

and with multiple MAR variables  $\{H_1, H_2, \dots, H_m\}$  this means

$$\Theta^{t+1} = \operatorname{argmax}_{\Theta} \left\{ \sum_{H_1} \sum_{H_2} \cdots \sum_{H_m} \alpha_H \log p(O, H | \Theta) \right\},$$

- In mixture models, EM **sums over all  $2^n$**  possible cluster assignments.
- In binary semi-supervised learning, EM **sums over all  $2^t$**  assignments to  $\tilde{y}$ .
- But **conditional independence** allows efficient calculation in the above cases.
  - The  $H$  are independent given  $\{O, \Theta\}$  which simplifies sums (see EM notes).
  - We'll cover general case when we discuss **probabilistic graphical models**.

## Today: Continuous-Latent Variables

- If  $H$  is continuous, the sums are replaced by **integrals**,

$$\log p(O|\Theta) = \log \left( \int_H p(O, H|\Theta) dH \right) \quad (\text{log-likelihood})$$

$$\Theta^{t+1} = \operatorname{argmax}_{\Theta} \left\{ \int_H \alpha_H \log p(O, H|\Theta) dH \right\} \quad (\text{EM update}),$$

where if we have 5 hidden variables  $\int_H$  means  $\int_{H_1} \int_{H_2} \int_{H_3} \int_{H_4} \int_{H_5}$ .

- Even with conditional independence these **might be hard**.

## Today: Continuous-Latent Variables

- If  $H$  is continuous, the sums are replaced by **integrals**,

$$\log p(O|\Theta) = \log \left( \int_H p(O, H|\Theta) dH \right) \quad (\text{log-likelihood})$$

$$\Theta^{t+1} = \operatorname{argmax}_{\Theta} \left\{ \int_H \alpha_H \log p(O, H|\Theta) dH \right\} \quad (\text{EM update}),$$

where if we have 5 hidden variables  $\int_H$  means  $\int_{H_1} \int_{H_2} \int_{H_3} \int_{H_4} \int_{H_5}$ .

- Even with conditional independence these **might be hard**.
- **Gaussian** assumptions allow efficient calculation of these integrals.
  - We'll cover general case when we get to discuss **Bayesian statistics**.



## Today: Continuous-Latent Variables

- In **mixture models**, we have a **discrete latent** variable  $z$ :
  - In mixture of Gaussians, if you know the cluster  $z$  then  $p(x|z)$  is a Gaussian.

## Today: Continuous-Latent Variables


- In **mixture models**, we have a **discrete latent** variable  $z$ :
  - In mixture of Gaussians, if you know the cluster  $z$  then  $p(x|z)$  is a Gaussian.
- In **latent-factor models**, we have **continuous latent** variables  $z$ :
  - In probabilistic PCA, if you know the latent-factors  $z$  then  $p(x|z)$  is a Gaussian.

## Today: Continuous-Latent Variables

- In **mixture models**, we have a **discrete latent** variable  $z$ :
  - In mixture of Gaussians, if you know the cluster  $z$  then  $p(x|z)$  is a Gaussian.
- In **latent-factor models**, we have **continuous latent** variables  $z$ :
  - In probabilistic PCA, if you know the latent-factors  $z$  then  $p(x|z)$  is a Gaussian.
- But what would a continuous  $z$  be useful for?
- Do we really need to start solving integrals?

## Today: Continuous-Latent Variables

- Data may live in a **low-dimensional manifold**:



<http://isomap.stanford.edu/handfig.html>

- **Mixtures are inefficient** at representing the 2D manifold.

## Principal Component Analysis (PCA)

- PCA replaces  $X$  with a lower-dimensional approximation  $Z$ .
  - Matrix  $Z$  has  $n$  rows, but typically far fewer columns.

## Principal Component Analysis (PCA)


- **PCA** replaces  $X$  with a lower-dimensional approximation  $Z$ .
  - Matrix  $Z$  has  $n$  rows, but typically **far fewer columns**.
- PCA is used for:
  - **Dimensionality reduction**: replace  $X$  with a lower-dimensional  $Z$ .

## Principal Component Analysis (PCA)

- **PCA** replaces  $X$  with a lower-dimensional approximation  $Z$ .
  - Matrix  $Z$  has  $n$  rows, but typically **far fewer columns**.
- PCA is used for:
  - **Dimensionality reduction**: replace  $X$  with a lower-dimensional  $Z$ .
  - **Outlier detection**: if PCA gives poor approximation of  $x^i$ , could be outlier.
  - **Basis for linear models**: use  $Z$  as features in regression model.

# Principal Component Analysis (PCA)

- **PCA** replaces  $X$  with a lower-dimensional approximation  $Z$ .
  - Matrix  $Z$  has  $n$  rows, but typically **far fewer columns**.
- PCA is used for:
  - **Dimensionality reduction**: replace  $X$  with a lower-dimensional  $Z$ .
  - **Outlier detection**: if PCA gives poor approximation of  $x^i$ , could be outlier.
  - **Basis for linear models**: use  $Z$  as features in regression model.
  - **Data visualization**: display  $z^i$  in a scatterplot.
  - **Factor discovering**: discover important hidden “factors” underlying data.





## PCA Notation

- PCA approximates the original matrix by factor-loadings  $Z$  and latent-factors  $W$ ,

$$X \approx ZW.$$

where  $Z \in \mathbb{R}^{n \times k}$ ,  $W \in \mathbb{R}^{k \times d}$ , and we assume columns of  $X$  have mean 0.

- We're trying to split redundancy in  $X$  into its important “parts”.

## PCA Notation

- PCA approximates the original matrix by factor-loadings  $Z$  and latent-factors  $W$ ,

$$X \approx ZW.$$

where  $Z \in \mathbb{R}^{n \times k}$ ,  $W \in \mathbb{R}^{k \times d}$ , and we assume columns of  $X$  have mean 0.

- We're trying to split redundancy in  $X$  into its important “parts”.
- We typically take  $k \ll d$  so this requires **far fewer parameters**:

$$\underbrace{\begin{bmatrix} \phantom{0} \\ \phantom{0} \\ \phantom{0} \\ \phantom{0} \\ \phantom{0} \end{bmatrix}}_{X \in \mathbb{R}^{n \times d}} \approx \underbrace{\begin{bmatrix} \phantom{0} \\ \phantom{0} \\ \phantom{0} \\ \phantom{0} \\ \phantom{0} \end{bmatrix}}_{Z \in \mathbb{R}^{n \times k}} \underbrace{\begin{bmatrix} \phantom{0} \\ \phantom{0} \\ \phantom{0} \\ \phantom{0} \\ \phantom{0} \end{bmatrix}}_{W \in \mathbb{R}^{k \times d}}$$

- Also computationally convenient:
  - $Xv$  costs  $O(nd)$  but  $Z(Wv)$  only costs  $O(nk + dk)$ .

## PCA Notation

- Using  $X \approx ZW$ , PCA approximates each examples  $x^i$  as

$$x^i \approx W^T z^i.$$

## PCA Notation

- Using  $X \approx ZW$ , PCA approximates each examples  $x^i$  as

$$x^i \approx W^T z^i.$$

- Usually we only need to estimate  $W$ :
  - If using least squares, then given  $W$  we can find  $z^i$  from  $x^i$  using

$$z^i = \underset{z}{\operatorname{argmin}} \|x^i - W^T z\|^2 = (WW^T)^{-1} W x^i.$$

## PCA Notation

- Using  $X \approx ZW$ , PCA approximates each examples  $x^i$  as

$$x^i \approx W^T z^i.$$

- Usually we only need to estimate  $W$ :
  - If using least squares, then given  $W$  we can find  $z^i$  from  $x^i$  using

$$z^i = \underset{z}{\operatorname{argmin}} \|x^i - W^T z\|^2 = (WW^T)^{-1} W x^i.$$

- We often assume that  $W^T$  is orthogonal:
  - This means that  $WW^T = I$ .
  - In this case we have  $z^i = W x^i$ .
- In standard formulations, solution only unique up to rotation:
  - Usually, we fit the rows of  $W$  sequentially for uniqueness.

## Two Classic Views on PCA

- PCA approximates the original matrix by latent-variables  $Z$  and latent-factors  $W$ ,

$$X \approx ZW.$$

where  $Z \in \mathbb{R}^{n \times k}$ ,  $W \in \mathbb{R}^{k \times d}$ .

- Two classical interpretations/derivations of PCA:

## Two Classic Views on PCA

- PCA approximates the original matrix by latent-variables  $Z$  and latent-factors  $W$ ,

$$X \approx ZW.$$

where  $Z \in \mathbb{R}^{n \times k}$ ,  $W \in \mathbb{R}^{k \times d}$ .

- Two classical interpretations/derivations of PCA:
  - ① Choose **latent-factors  $W$  to minimize error** (“synthesis view”):

$$\operatorname{argmin}_{Z \in \mathbb{R}^{n \times k}, W \in \mathbb{R}^{k \times d}} \|X - ZW\|_F^2 = \sum_{i=1}^n \sum_{j=1}^d (x_j^i - (w_j)^T z^i)^2.$$

## Two Classic Views on PCA

- PCA approximates the original matrix by latent-variables  $Z$  and latent-factors  $W$ ,

$$X \approx ZW.$$

where  $Z \in \mathbb{R}^{n \times k}$ ,  $W \in \mathbb{R}^{k \times d}$ .

- Two classical interpretations/derivations of PCA:

- 1 Choose **latent-factors  $W$  to minimize error** (“synthesis view”):


$$\operatorname{argmin}_{Z \in \mathbb{R}^{n \times k}, W \in \mathbb{R}^{k \times d}} \|X - ZW\|_F^2 = \sum_{i=1}^n \sum_{j=1}^d (x_j^i - (w_j)^T z^i)^2.$$

- 2 Choose **orthogonal latent-factors  $W$  to maximize variance** (“analysis view”):


$$\begin{aligned} \operatorname{argmax}_{W \in \mathbb{R}^{k \times d}} &= \sum_{i=1}^n \|z^i - \mu_z\|^2 = \sum_{i=1}^n \|Wx^i\|^2 && (z^i = Wx^i \text{ and } \mu_z = 0) \\ &= \operatorname{Tr}(WX^T XW^T) = n \operatorname{Tr}(W^T W S) && (\text{where } S \text{ is sample covariance}) \end{aligned}$$




## PCA in One Dimension




## PCA in One Dimension




## PCA in One Dimension



## PCA in One Dimension



## PCA in One Dimension



## Probabilistic PCA

- With zero-mean (“centered”) data, in PCA we assume that

$$x \approx W^T z.$$

- In **probabilistic PCA** we assume that

$$x \sim \mathcal{N}(W^T z, \sigma^2 I), \quad z \sim \mathcal{N}(0, I).$$

(we can actually use any Gaussian density for  $z$ )

## Probabilistic PCA

- With zero-mean (“centered”) data, in PCA we assume that

$$x \approx W^T z.$$

- In **probabilistic PCA** we assume that

$$x \sim \mathcal{N}(W^T z, \sigma^2 I), \quad z \sim \mathcal{N}(0, I).$$

(we can actually use any Gaussian density for  $z$ )

- Since  $z$  is hidden, our observed likelihood integrates over  $z$ ,

$$p(x|W) = \int_z p(x, z|W) dz.$$

- Looks ugly, but can be computed due to the Gaussians assumptions:
  - This marginal distribution is Gaussian.

## Manipulating Gaussians

- From the assumptions of the previous slide we have

$$p(x|z, W) \propto \exp\left(-\frac{(x - W^T z)^T (x - W^T z)}{2\sigma^2}\right), \quad p(z) \propto \exp\left(-\frac{z^T z}{2}\right).$$



## Manipulating Gaussians

- From the assumptions of the previous slide we have

$$p(x|z, W) \propto \exp\left(-\frac{(x - W^T z)^T (x - W^T z)}{2\sigma^2}\right), \quad p(z) \propto \exp\left(-\frac{z^T z}{2}\right).$$

- Multiplying and expanding we get

$$p(x, z|W) = p(x|z, W)p(z|W)$$

## Manipulating Gaussians

- From the assumptions of the previous slide we have

$$p(x|z, W) \propto \exp\left(-\frac{(x - W^T z)^T (x - W^T z)}{2\sigma^2}\right), \quad p(z) \propto \exp\left(-\frac{z^T z}{2}\right).$$

- Multiplying and expanding we get

$$\begin{aligned} p(x, z|W) &= p(x|z, W)p(z|W) \\ &= p(x|z, W)p(z) \end{aligned} \quad (z \perp W)$$

## Manipulating Gaussians

- From the assumptions of the previous slide we have

$$p(x|z, W) \propto \exp\left(-\frac{(x - W^T z)^T (x - W^T z)}{2\sigma^2}\right), \quad p(z) \propto \exp\left(-\frac{z^T z}{2}\right).$$

- Multiplying and expanding we get

$$\begin{aligned} p(x, z|W) &= p(x|z, W)p(z|W) \\ &= p(x|z, W)p(z) && (z \perp W) \\ &\propto \exp\left(-\frac{(x - W^T z)^T (x - W^T z)}{2\sigma^2} - \frac{z^T z}{2}\right) \end{aligned}$$

## Manipulating Gaussians

- From the assumptions of the previous slide we have

$$p(x|z, W) \propto \exp\left(-\frac{(x - W^T z)^T (x - W^T z)}{2\sigma^2}\right), \quad p(z) \propto \exp\left(-\frac{z^T z}{2}\right).$$

- Multiplying and expanding we get

$$\begin{aligned} p(x, z|W) &= p(x|z, W)p(z|W) \\ &= p(x|z, W)p(z) && (z \perp W) \\ &\propto \exp\left(-\frac{(x - W^T z)^T (x - W^T z)}{2\sigma^2} - \frac{z^T z}{2}\right) \\ &= \exp\left(-\frac{x^T x - x^T W^T z - z^T W x + z^T W W^T z}{2\sigma^2} + \frac{z^T z}{2}\right) \end{aligned}$$

## Manipulating Gaussians

- So the “complete” likelihood satisfies

$$p(x, z|W) \propto \exp \left( -\frac{x^T x - x^T W^T z - z^T W x + z^T W W^T z}{2\sigma^2} + \frac{z^T z}{2} \right)$$

## Manipulating Gaussians

- So the “complete” likelihood satisfies

$$\begin{aligned} p(x, z|W) &\propto \exp \left( -\frac{x^T x - x^T W^T z - z^T W x + z^T W W^T z}{2\sigma^2} + \frac{z^T z}{2} \right) \\ &= \exp \left( -\frac{1}{2} \left( x^T \left( \frac{1}{\sigma^2} I \right) x + x^T \left( \frac{1}{\sigma^2} W^T \right) z + z^T \left( \frac{1}{\sigma^2} W \right) x + z^T \left( \frac{1}{\sigma^2} W W^T + I \right) z \right) \right), \end{aligned}$$

## Manipulating Gaussians

- So the “complete” likelihood satisfies

$$\begin{aligned} p(x, z|W) &\propto \exp\left(-\frac{x^T x - x^T W^T z - z^T W x + z^T W W^T z}{2\sigma^2} + \frac{z^T z}{2}\right) \\ &= \exp\left(-\frac{1}{2}\left(x^T\left(\frac{1}{\sigma^2}I\right)x + x^T\left(\frac{1}{\sigma^2}W^T\right)z + z^T\left(\frac{1}{\sigma^2}W\right)x + z^T\left(\frac{1}{\sigma^2}WW^T + I\right)z\right)\right), \end{aligned}$$

- We can re-write the exponent as a quadratic form,

$$p(x, z|W) \propto \exp\left(-\frac{1}{2}\begin{bmatrix} z^T & x^T \end{bmatrix} \begin{bmatrix} \frac{1}{\sigma^2}WW^T + I & -\frac{1}{\sigma^2}W \\ -\frac{1}{\sigma^2}W^T & \frac{1}{\sigma^2}I \end{bmatrix} \begin{bmatrix} z \\ x \end{bmatrix}\right),$$

## Manipulating Gaussians

- So the “complete” likelihood satisfies

$$\begin{aligned} p(x, z|W) &\propto \exp \left( -\frac{x^T x - x^T W^T z - z^T W x + z^T W W^T z}{2\sigma^2} + \frac{z^T z}{2} \right) \\ &= \exp \left( -\frac{1}{2} \left( x^T \left( \frac{1}{\sigma^2} I \right) x + x^T \left( \frac{1}{\sigma^2} W^T \right) z + z^T \left( \frac{1}{\sigma^2} W \right) x + z^T \left( \frac{1}{\sigma^2} W W^T + I \right) z \right) \right), \end{aligned}$$

- We can re-write the exponent as a quadratic form,

$$p(x, z|W) \propto \exp \left( -\frac{1}{2} \begin{bmatrix} z^T & x^T \end{bmatrix} \begin{bmatrix} \frac{1}{\sigma^2} W W^T + I & -\frac{1}{\sigma^2} W \\ -\frac{1}{\sigma^2} W^T & \frac{1}{\sigma^2} I \end{bmatrix} \begin{bmatrix} z \\ x \end{bmatrix} \right),$$

- This has the form of a Gaussian distribution,

$$p(v|W) \propto \exp \left( -\frac{1}{2} v^T \Sigma^{-1} v \right),$$

with  $v = \begin{bmatrix} z \\ x \end{bmatrix}$ ,  $\mu = 0$ , and  $\Sigma^{-1} = \begin{bmatrix} \frac{1}{\sigma^2} W W^T + I & -\frac{1}{\sigma^2} W \\ -\frac{1}{\sigma^2} W^T & \frac{1}{\sigma^2} I \end{bmatrix}$ .



## Manipulating Gaussians

- Remember that if we write multivariate Gaussian in partitioned form,

$$\begin{bmatrix} x \\ z \end{bmatrix} \sim \mathcal{N} \left( \begin{bmatrix} \mu_x \\ \mu_z \end{bmatrix}, \begin{bmatrix} \Sigma_{xx} & \Sigma_{xz} \\ \Sigma_{zx} & \Sigma_{zz} \end{bmatrix} \right),$$

then the marginal distribution  $p(x)$  (integrating over  $z$ ) is given by

$$x \sim \mathcal{N}(\mu_x, \Sigma_{xx}).$$


## Manipulating Gaussians

- Remember that if we write multivariate Gaussian in partitioned form,

$$\begin{bmatrix} x \\ z \end{bmatrix} \sim \mathcal{N} \left( \begin{bmatrix} \mu_x \\ \mu_z \end{bmatrix}, \begin{bmatrix} \Sigma_{xx} & \Sigma_{xz} \\ \Sigma_{zx} & \Sigma_{zz} \end{bmatrix} \right),$$

then the marginal distribution  $p(x)$  (integrating over  $z$ ) is given by

$$x \sim \mathcal{N}(\mu_x, \Sigma_{xx}).$$



## Manipulating Gaussians

- Remember that if we write multivariate Gaussian in partitioned form,

$$\begin{bmatrix} x \\ z \end{bmatrix} \sim \mathcal{N} \left( \begin{bmatrix} \mu_x \\ \mu_z \end{bmatrix}, \begin{bmatrix} \Sigma_{xx} & \Sigma_{xz} \\ \Sigma_{zx} & \Sigma_{zz} \end{bmatrix} \right),$$

then the marginal distribution  $p(x)$  (integrating over  $z$ ) is given by

$$x \sim \mathcal{N}(\mu_x, \Sigma_{xx}).$$

- For probabilistic PCA we assume  $\mu_x = 0$ , but we partitioned  $\Sigma^{-1}$  instead of  $\Sigma$ .

## Manipulating Gaussians

- Remember that if we write multivariate Gaussian in partitioned form,

$$\begin{bmatrix} x \\ z \end{bmatrix} \sim \mathcal{N} \left( \begin{bmatrix} \mu_x \\ \mu_z \end{bmatrix}, \begin{bmatrix} \Sigma_{xx} & \Sigma_{xz} \\ \Sigma_{zx} & \Sigma_{zz} \end{bmatrix} \right),$$

then the marginal distribution  $p(x)$  (integrating over  $z$ ) is given by

$$x \sim \mathcal{N}(\mu_x, \Sigma_{xx}).$$

- For probabilistic PCA we assume  $\mu_x = 0$ , but we partitioned  $\Sigma^{-1}$  instead of  $\Sigma$ .
- To get  $\Sigma$  we can use a [matrix inversion lemma](#),

$$\Sigma = \begin{bmatrix} \frac{1}{\sigma^2} W W^T + I & -\frac{1}{\sigma^2} W \\ -\frac{1}{\sigma^2} W^T & \frac{1}{\sigma^2} I \end{bmatrix}^{-1} = \begin{bmatrix} W^T W + \sigma^2 I & W^T \\ W & I \end{bmatrix},$$

which gives that [solution to integrating over  \$z\$](#)  is

$$x|W \sim \mathcal{N}(0, W^T W + \sigma^2 I).$$

## Notes on Probabilistic PCA

- Negative log-likelihood of observed data has the form

$$-\log p(x|W) = \frac{n}{2} \text{Tr}(SC) + \frac{n}{2} \log |C| + \text{const.},$$

where  $C = W^T W + \sigma^2 I$  and  $S = X^T X$ .

## Notes on Probabilistic PCA

- Negative log-likelihood of observed data has the form

$$-\log p(x|W) = \frac{n}{2} \text{Tr}(SC) + \frac{n}{2} \log |C| + \text{const.},$$

where  $C = W^T W + \sigma^2 I$  and  $S = X^T X$ .

- Not convex, but **non-global stationary points are saddle points**.
- Regular PCA is obtained as limit of  $\sigma \rightarrow 0$  (as  $\log |C| \rightarrow 0$ ).

## Notes on Probabilistic PCA

- Negative log-likelihood of observed data has the form

$$-\log p(x|W) = \frac{n}{2} \text{Tr}(SC) + \frac{n}{2} \log |C| + \text{const.},$$

where  $C = W^T W + \sigma^2 I$  and  $S = X^T X$ .

- Not convex, but **non-global stationary points are saddle points**.
- Regular PCA is obtained as limit of  $\sigma \rightarrow 0$  (as  $\log |C| \rightarrow 0$ ).
- Can **reduce cost from  $O(d^3)$  to  $O(k^3)$**  with matrix inversion/determinant lemmas:
  - Allows us to work with  $WW^T$  instead of  $W^T W$ .
- We can get  $p(z|x, W)$  using that conditional of Gaussians is Gaussian.

## Notes on Probabilistic PCA

- Negative log-likelihood of observed data has the form

$$-\log p(x|W) = \frac{n}{2} \text{Tr}(SC) + \frac{n}{2} \log |C| + \text{const.},$$

where  $C = W^T W + \sigma^2 I$  and  $S = X^T X$ .

- Not convex, but **non-global stationary points are saddle points**.
- Regular PCA is obtained as limit of  $\sigma \rightarrow 0$  (as  $\log |C| \rightarrow 0$ ).
- Can **reduce cost from  $O(d^3)$  to  $O(k^3)$**  with matrix inversion/determinant lemmas:
  - Allows us to work with  $WW^T$  instead of  $W^T W$ .
- We can get  $p(z|x, W)$  using that conditional of Gaussians is Gaussian.
- We could consider different distribution for  $x^i|z^i$  (but integrals are ugly):
  - E.g., Laplace of student if you want it to be robust.
  - E.g., logistic or softmax if you have discrete  $x_j^i$ .



## Generalizations of Probabilistic PCA

- Why do we need a probabilistic interpretation of PCA?
  - Good excuse to play with Gaussian identities and matrix formulas?

## Generalizations of Probabilistic PCA

- Why do we need a probabilistic interpretation of PCA?
  - Good excuse to play with Gaussian identities and matrix formulas?
- We now understand that **PCA fits a Gaussian with restricted covariance**:
  - Hope is that  $W^T W + \sigma I$  is a good approximation of full covariance  $X^T X$ .

## Generalizations of Probabilistic PCA

- Why do we need a probabilistic interpretation of PCA?
  - Good excuse to play with Gaussian identities and matrix formulas?
- We now understand that **PCA fits a Gaussian with restricted covariance**:
  - Hope is that  $W^T W + \sigma I$  is a good approximation of full covariance  $X^T X$ .
  - We can do fancy things like **mixtures of PCA** models.




Figure 8: Comparison of an 8-component diagonal variance Gaussian mixture model with a mixture of PPCA model. The upper two plots give a view perpendicular to the major

<http://www.miketipping.com/papers/met-mppca.pdf>

- Lets us understand connection between PCA and **factor analysis**.

# Factor Analysis

- **Factor analysis** (FA) is a method for discovering latent-factors.
- Historical applications are measures of intelligence and personality traits.
  - Some controversy, like trying to find factors of intelligence due to race.  
(without normalizing for socioeconomic factors)

Trait	Description
<b>O</b> penness	Being curious, original, intellectual, creative, and open to new ideas.
<b>C</b> onscientiousness	Being organized, systematic, punctual, achievement-oriented, and dependable.
<b>E</b> xtraversion	Being outgoing, talkative, sociable, and enjoying social situations.
<b>A</b> greeableness	Being affable, tolerant, sensitive, trusting, kind, and warm.
<b>N</b> euroticism	Being anxious, irritable, temperamental, and moody.

<https://new.edu/resources/big-5-personality-traits>

- But a standard tool and widely-used across science and engineering.

# Factor Analysis

- FA approximates the original matrix by latent-variables  $Z$  and latent-factors  $W$ ,

$$X \approx ZW.$$


- Which should sound familiar...

# Factor Analysis

- FA approximates the original matrix by latent-variables  $Z$  and latent-factors  $W$ ,


$$X \approx ZW.$$

- Which should sound familiar...
- Are PCA and FA the same?
  - Both are more than 100 years old.
  - People are still fighting about whether they are the same:
    - Doesn't help that some software packages run PCA when you call FA.

Google  

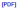
[All](#) [Images](#) [Videos](#) [News](#) [Maps](#) [More ▾](#) [Search tools](#)

About 358,000 results (0.17 seconds)


 **Principal Component Analysis versus Exploratory Factor ...**  
[www2.sas.com/proceedings/sug30/203-30.pdf](http://www2.sas.com/proceedings/sug30/203-30.pdf) ▸  
 by DD Suhr - Cited by 118 - Related articles  
 1. Paper 203-30. Principal Component Analysis vs. Exploratory Factor Analysis.  
 Diana D. Suhr, Ph.D. University of Northern Colorado. Abstract. Principal ...

**pca - What are the differences between Factor Analysis and ...**  
[stats.stackexchange.com/.../what-are-the-differences-between-factor-anal...](http://stats.stackexchange.com/.../what-are-the-differences-between-factor-anal...) ▸  
 Aug 12, 2010 - Principal Component Analysis (PCA) and Common Factor Analysis (CFA) ... differently one has to interpret the strength of loadings in PCA vs.


**What are the differences between principal components ...**  
[support.minitab.com/.../factor-analysis/differences-between-pca-and-factor...](http://support.minitab.com/.../factor-analysis/differences-between-pca-and-factor...) ▸  
 Principal Components Analysis and Factor Analysis are similar because both procedures are used to simplify the structure of a set of variables. However, the ...

 **Principal Components Analysis - UNT**  
<https://www.unt.edu/rss/class/.../Principal%20Components%20Analysis.p...> ▸  
 PCA vs. Factor Analysis • It is easy to make the mistake in assuming that these are the same techniques, though in some ways exploratory factor analysis and ...


**Factor analysis versus Principal Components Analysis (PCA)**  
[psych.wisc.edu/henriques/pca.html](http://psych.wisc.edu/henriques/pca.html) ▸  
 Jun 19, 2010 - Factor analysis versus PCA. These techniques are typically used to analyze groups of correlated variables representing one or more common ...

 **Principal Component Analysis and Factor Analysis**  
[www.stats.ox.ac.uk/~ripley/MultAnal\\_1RT2007/PC-FA.pdf](http://www.stats.ox.ac.uk/~ripley/MultAnal_1RT2007/PC-FA.pdf) ▸  
 where D is diagonal with non-negative and decreasing values and U and V ...  
 Factor analysis and PCA are often confused, and indeed SPSS has PCA as.

**How can I decide between using principal components ...**  
[https://www.researchgate.net/.../How\\_can\\_I\\_decide\\_between\\_using\\_prin...](https://www.researchgate.net/.../How_can_I_decide_between_using_prin...) ▸  
 Factor analysis (FA) is a group of statistical methods used to understand and simplify patterns ... Retrieved from <http://pareonline.net/getn.asp?v=10&n=7> ...  
 Principal component analysis (PCA) is a method of factor extraction (the second step ...

 **Exploratory Factor Analysis and Principal Component An...**  
[www.lesahoffman.com/948/948\\_Lecture2\\_EFA\\_PCA.pdf](http://www.lesahoffman.com/948/948_Lecture2_EFA_PCA.pdf) ▸  
 2 very different schools of thought on exploratory factor analysis (EFA) vs. principal components analysis (PCA) > EFA and PCA are TWO ENTIRELY ...

**Factor analysis - Wikipedia, the free encyclopedia**  
[https://en.wikipedia.org/wiki/Factor\\_analysis](https://en.wikipedia.org/wiki/Factor_analysis) ▸  
 Jump to: [Exploratory factor analysis versus principal components](#) ... [edit]. See also: Principal component analysis and Exploratory factor analysis.

 **The Truth about PCA and Factor Analysis**  
[www.stat.cmu.edu/~cshalizi/350/lectures/13/lecture-13.pdf](http://www.stat.cmu.edu/~cshalizi/350/lectures/13/lecture-13.pdf) ▸  
 Sep 28, 2009 - nents and factor analysis, we'll wrap up by looking at their uses and

## PCA vs. Factor Analysis

- In probabilistic PCA we assume

$$x|z \sim \mathcal{N}(W^T z, \sigma^2 I), \quad z \sim \mathcal{N}(0, I),$$

and we obtain PCA as  $\sigma \rightarrow 0$ .



## PCA vs. Factor Analysis

- In probabilistic PCA we assume

$$x|z \sim \mathcal{N}(W^T z, \sigma^2 I), \quad z \sim \mathcal{N}(0, I),$$

and we obtain PCA as  $\sigma \rightarrow 0$ .

- In FA we assume

$$x|z \sim \mathcal{N}(W^T z, D), \quad z \sim \mathcal{N}(0, I),$$

where  $D$  is a diagonal matrix.

## PCA vs. Factor Analysis

- In probabilistic PCA we assume

$$x|z \sim \mathcal{N}(W^T z, \sigma^2 I), \quad z \sim \mathcal{N}(0, I),$$

and we obtain PCA as  $\sigma \rightarrow 0$ .

- In FA we assume

$$x|z \sim \mathcal{N}(W^T z, D), \quad z \sim \mathcal{N}(0, I),$$

where  $D$  is a diagonal matrix.

- The difference is that you can have a **noise variance for each dimension**.
- Repeating the previous exercise we get that

$$x \sim \mathcal{N}(0, W^T W + D).$$

## PCA vs. Factor Analysis

- We can write non-centered versions of both models:
  - Probabilistic PCA:

$$x|z \sim \mathcal{N}(W^T z + \mu, \sigma^2 I), \quad z \sim \mathcal{N}(0, I),$$

- Factor analysis:

$$x|z \sim \mathcal{N}(W^T z + \mu, D), \quad z \sim \mathcal{N}(0, I),$$

where  $D$  is a diagonal matrix.

## PCA vs. Factor Analysis

- We can write non-centered versions of both models:

- Probabilistic PCA:

$$x|z \sim \mathcal{N}(W^T z + \mu, \sigma^2 I), \quad z \sim \mathcal{N}(0, I),$$

- Factor analysis:

$$x|z \sim \mathcal{N}(W^T z + \mu, D), \quad z \sim \mathcal{N}(0, I),$$

where  $D$  is a diagonal matrix.



- A different perspective is that these models assume

$$x = W^T z + \epsilon,$$

where PPCA has  $\epsilon \sim \mathcal{N}(\mu, \sigma^2 I)$  and FA has  $\epsilon \sim \mathcal{N}(\mu, D)$ .

- So in FA has extra degrees of freedom in variance of individual variables.

# PCA vs. Factor Analysis



[http:](http://stats.stackexchange.com/questions/1576/what-are-the-differences-between-factor-analysis-and-principal-component-analysis)

[//stats.stackexchange.com/questions/1576/what-are-the-differences-between-factor-analysis-and-principal-component-analysis](http://stats.stackexchange.com/questions/1576/what-are-the-differences-between-factor-analysis-and-principal-component-analysis)

Remember in 340 that difference with PCA and ISOMAP/t-SNE was huge.

## Factor Analysis Discussion

- No closed-form solution for FA, and can find different local optima.
- Unlike PCA, FA doesn't change if you scale variables.
  - FA doesn't chase large-noise features that are uncorrelated with other features.

## Factor Analysis Discussion

- No closed-form solution for FA, and can find different local optima.
- Unlike PCA, FA doesn't change if you scale variables.
  - FA doesn't chase large-noise features that are uncorrelated with other features.
- Unlike PCA, FA changes if you rotate data.
- Similar to PCA, objective only depends on  $W^T W$  so you can rotate/mirror the factors

$$W^T W = W^T \underbrace{Q^T Q}_I W = (WQ)^T (WQ),$$

for an orthogonal matrix  $Q$ .

- So you **can't interpret multiple factors as being unique**.

# Summary

- **Monotonicity of EM:** EM is guaranteed not to decrease likelihood.



# Summary

- **Monotonicity of EM:** EM is guaranteed not to decrease likelihood.
- **Kernel density estimation:** Non-parametric continuous density estimation method.

## Summary

- **Monotonicity of EM**: EM is guaranteed not to decrease likelihood.
- **Kernel density estimation**: Non-parametric continuous density estimation method.
- **PCA** is a classic method for dimensionality reduction.

## Summary

- **Monotonicity of EM**: EM is guaranteed not to decrease likelihood.
- **Kernel density estimation**: Non-parametric continuous density estimation method.
- **PCA** is a classic method for dimensionality reduction.
- **Probabilistic PCA** is a continuous latent-variable probabilistic generalization.

## Summary

- **Monotonicity of EM**: EM is guaranteed not to decrease likelihood.
- **Kernel density estimation**: Non-parametric continuous density estimation method.
- **PCA** is a classic method for dimensionality reduction.
- **Probabilistic PCA** is a continuous latent-variable probabilistic generalization.
- **Factor analysis** extends probabilistic PCA with different noise in each dimension.
- Next time: the algorithm we didn't cover in 340 from the list of  
"The 10 Algorithms Machine Learning Engineers Need to Know".

## Bonus Slide: Mixture of Experts

- Classic generative model for supervised learning uses

$$p(y^i|x^i) \propto p(x^i|y^i)p(y^i),$$

and typically  $p(x^i|y^i)$  is assumed by Gaussian (LDA) or independent (naive Bayes).

## Bonus Slide: Mixture of Experts

- Classic generative model for supervised learning uses

$$p(y^i|x^i) \propto p(x^i|y^i)p(y^i),$$

and typically  $p(x^i|y^i)$  is assumed by Gaussian (LDA) or independent (naive Bayes).

- But we could allow more flexibility by using a mixture model,

$$p(x^i|y^i) = \sum_{c=1}^k p(z^i = c|y^i)p(x^i|z^i = c, y^i).$$

## Bonus Slide: Mixture of Experts

- Classic generative model for supervised learning uses

$$p(y^i|x^i) \propto p(x^i|y^i)p(y^i),$$

and typically  $p(x^i|y^i)$  is assumed by Gaussian (LDA) or independent (naive Bayes).

- But we could allow more flexibility by using a mixture model,

$$p(x^i|y^i) = \sum_{c=1}^k p(z^i = c|y^i)p(x^i|z^i = c, y^i).$$

- Instead of a generative model, we could also take a mixture of regression models,

$$p(y^i|x^i) = \sum_{c=1}^k p(z^i = c|x^i)p(y^i|z^i = c, x^i).$$

- Called a “mixture of experts” model:
  - Each regression model is an “expert” for certain values of  $x^i$ .