

CPSC 540 Assignment 4 (due March 20)

Graphical Models and Paper Review

1 Markov Chains

1.1 Sampling, Inference, and Decoding

The function `example_markovChain.m` loads the initial state probabilities and transition probabilities for three Markov chain models on d binary variables,

$$p(x_1, x_2, \dots, x_d) = p(x_1) \prod_{j=2}^d p(x_j | x_{j-1}).$$

It then tries to find the optimal decoding (the most likely assignment to the variables $\{x_1, x_2, \dots, x_d\}$) in each of the three chains. In the demo, decoding is done by enumerating all possible assignments to the variables. This works for the first two chains as they only have 4 variables, but is too slow on the last chain because it has 30 variables. In this question you'll explore two ways to estimate the marginals in third Markov chain and two ways to estimate the most-probable sequence.

1. Write a function, `sampleAncestral.m`, that uses ancestral sampling to sample sequence x . [Hand in this code and report all the univariate marginal probabilities using a Monte Carlo estimate based on 10000 samples.](#)
2. Write a function, `marginalCK.m`, that uses the CK equations to compute the exact univariate marginals. [Hand in this code and report all exact univariate marginals.](#)
3. Write a function, `marginalDecode.m`, that returns the sequence of states x_j that maximize the marginal probability $p(x_j)$ (for each j). [Hand in this code and report the sequence of most likely states.](#)
4. Write a function, `viterbiDecode.m`, that implements the Viterbi decoding algorithm for Markov chains. [Hand in this code and report the optimal decoding of the third Markov chain.](#)

Hint: for parts 2-4, you can use a 2 by d matrix M to represent the dynamic programming table, and for part 4 you can use another matrix B containing the argmax values that lead to each entry in the table.

Solution

1 and 2.

Using the script `sampleAncestral.m` along with the script `sampleDiscrete.m` to produce k samples in $\mathcal{O}(\log(k))$ we produced a 10000 from the Markov Chain with 30 elements and we also calculated the exact marginals using the script `marginalsCK.m`. To run the scripts we used `marginalsCalculation.m` shown below.

```
1 load viterbiData.mat
2
```

```

3
4 n=10000; %Number of samples
5
6 X=sampleAncestral(p0,pT_long,n);
7 marginals=sum(X==1)/n;
8 marginals=[marginals;1-marginals];
9 display('The marginals obtained using Monte Carlo estimation: ')
10 marginals '
11
12 marginals2=marginalCK(p0,pT_long);
13
14 display('The marginals obtained using CK equations: ')
15 marginals2 '

```

The output obtained by running the above script is the following

```
>>marginalCalculation
```

The marginals obtained using Monte Carlo estimation:

```
ans =
```

0.5999	0.4001
0.4720	0.5280
0.9035	0.0965
0.9704	0.0296
0.7755	0.2245
0.2263	0.7737
0.0533	0.9467
0.3323	0.6677
0.3163	0.6837
0.6032	0.3968
0.3505	0.6495
0.7586	0.2414
0.3695	0.6305
0.1321	0.8679
0.3646	0.6354
0.4015	0.5985
0.7233	0.2767
0.7879	0.2121
0.3122	0.6878
0.6750	0.3250
0.5848	0.4152
0.2511	0.7489
0.4032	0.5968
0.2293	0.7707
0.7452	0.2548
0.6203	0.3797
0.6039	0.3961
0.7432	0.2568
0.4740	0.5260
0.4347	0.5653

The marginals obtained using CK equations:

ans =

0.6000	0.4000
0.4681	0.5319
0.8998	0.1002
0.9718	0.0282
0.7779	0.2221
0.2267	0.7733
0.0542	0.9458
0.3359	0.6641
0.3204	0.6796
0.5902	0.4098
0.3601	0.6399
0.7479	0.2521
0.3726	0.6274
0.1313	0.8687
0.3744	0.6256
0.4103	0.5897
0.7240	0.2760
0.7877	0.2123
0.3147	0.6853
0.6698	0.3302
0.5811	0.4189
0.2541	0.7459
0.4037	0.5963
0.2403	0.7597
0.7378	0.2622
0.6211	0.3789
0.6060	0.3940
0.7419	0.2581
0.4644	0.5356
0.4380	0.5620

As we can see convergence is slow, since the approximation with 10000 samples gives accurate result to just one significant figure.

3.

Now we proceed to maximize the marginals, i.e. we calculate

$$\max_{x_1, \dots, x_{30}} \mathbb{P}(x_1) \dots \mathbb{P}(x_{30}).$$

To calculate this quantity we used the script marginalDecode.m. The output obtained shows the most likely chain given the maximization problem stated above

```
>>marginalDecode(p0,pT_long)
```

ans =

Columns 1 through 13

1 2 1 1 1 2 2 2 2 1 2 1 2

Columns 14 through 26

2 2 2 1 1 2 1 1 2 2 2 1 1

Columns 27 through 30

1 1 2 2

The probability to obtain this chain can be easily calculated using

$$\mathbb{P}(1, 2, 1, \dots, 2, 2) = \mathbb{P}(x_1 = 1)\mathbb{P}(x_2 = 2|x_1 = 1)\mathbb{P}(x_3 = 1|x_1 = 1) \cdots \mathbb{P}(x_{30} = 2|x_{29} = 2).$$

The result obtained using the probability transition matrices contained in *pT_long(:, :, 1 : 29)* is

$$\mathbb{P}(1, 2, 1, \dots, 2, 2) = 5.0639 \times 10^{-5}.$$

In the next section we will see how this compares with the probability for the chain with highest probability.
4.

Now we calculate the chain with maximum probability using the script *viterbiDecode.m*. The results are shown below

```
>>viterbiDecode(p0,pT_long)
```

```
ans =
```

Columns 1 through 13

1 1 1 1 1 2 2 2 2 1 2 1 2

Columns 14 through 26

2 2 2 2 1 2 1 1 2 2 2 1 1

Columns 27 through 30

1 1 1 1

The probability to get this chain is given by

$$\mathbb{P}(1, 1, \dots, 1, 1) = 1.7293 \times 10^{-4}.$$

This means that getting the chain with the maximum probability is almost 3.5 times more likely to see it than the chain that is obtained by maximizing the marginal probabilities.

1.2 Conditioning

The long sequence from the previous question usually starts with state 1 and most of the time ends in state 2. In this question you'll consider conditioning on these events not happening. First, compute the following quantities which can be done using your functions from the previous question:

1. Report all the univariate conditional probabilities $p(x_j|x_1 = 2)$ obtained using a Monte Carlo estimate based on 10000 samples.

2. Report all the exact univariate conditionals $p(x_j|x_1 = 2)$.
3. Report the sequence beginning with $x_1 = 2$ that has the highest probability.
4. Report the sequence ending with $x_d = 1$ that has the highest probability.

Hint: these conditions can be done by changing the input to the functions from the previous question.

Next consider the following cases (which require implementing an extra rejection step or backward phase):

5. Report all the univariate conditional probabilities $p(x_j|x_d = 1)$ obtained using a Monte Carlo estimate based on 10000 samples and rejection sampling. Also report the number of samples accepted among the 10000 samples.
6. Write a function, *sampleBackwards.m* that uses backwards sampling to sample sequences where $x_d = 1$. Hand in this code and report all the univariate conditional probabilities $p(x_j|x_d = 1)$ obtained using a Monte Carlo estimate based on 10000 samples.
7. Write a function, *forwardBackwards.m* that is able compute all exact univariate conditionals $p(x_j|x_d = 1)$ in $O(dk^2)$. Hand in the code and report all the exact univariate conditionals $p(x_j|x_d = 1)$.

Solutions

1 and 2

Using the script *sampleAncestral.m* and using the initial probability p_0 as $p_0 = [1, 0]$ we produce a 10000 samples from the Model where every chain starts at $x_1 = 1$. With these samples we do a Monte Carlo integration to approximate the marginals. Using the same initial probability p_0 we use the script *marginalCK.m* to calculate the exact conditional marginals. Below we show the results obtained from the Monte Carlo approximation and the exact calculation of the marginals, where as before the j -th row in the first column is $\mathbb{P}(x_j = 1|x_1 = 2)$ and the second column is $\mathbb{P}(x_j = 2|x_1 = 2)$

```
>>question12
```

The conditional probability $p(x_j|x_1=2)$ is given by:

```
ans =
```

0	1.0000
0.0619	0.9381
0.9327	0.0673
0.9747	0.0253
0.7754	0.2246
0.2221	0.7779
0.0600	0.9400
0.3428	0.6572
0.3142	0.6858
0.5965	0.4035
0.3562	0.6438
0.7500	0.2500
0.3777	0.6223
0.1287	0.8713
0.3767	0.6233
0.4081	0.5919
0.7273	0.2727
0.7871	0.2129
0.3094	0.6906

0.6746	0.3254
0.5877	0.4123
0.2649	0.7351
0.4004	0.5996
0.2433	0.7567
0.7336	0.2664
0.6164	0.3836
0.6015	0.3985
0.7438	0.2562
0.4595	0.5405
0.4333	0.5667

The exact univariate conditionals are:

M =

0	1.0000
0.0634	0.9366
0.9297	0.0703
0.9756	0.0244
0.7789	0.2211
0.2265	0.7735
0.0542	0.9458
0.3359	0.6641
0.3204	0.6796
0.5902	0.4098
0.3601	0.6399
0.7479	0.2521
0.3726	0.6274
0.1313	0.8687
0.3744	0.6256
0.4103	0.5897
0.7240	0.2760
0.7877	0.2123
0.3147	0.6853
0.6698	0.3302
0.5811	0.4189
0.2541	0.7459
0.4037	0.5963
0.2403	0.7597
0.7378	0.2622
0.6211	0.3789
0.6060	0.3940
0.7419	0.2581
0.4644	0.5356
0.4380	0.5620

As before we can see that the Monte Carlo integration with 10000 sample only gives accuracy of one significant digit.

3.

Now we use the script viterbiDecoding using the same initial probability p_0 as above, with this we get

>>question12

The sequence with highest probability starting with $x_1=2$ is:

ans =

Columns 1 through 13

2 2 1 1 1 2 2 2 2 1 2 1 2

Columns 14 through 26

2 2 2 2 1 2 1 1 2 2 2 1 1

Columns 27 through 30

1 1 1 1

4.

To calculate the sequence with highest probability starting with $x_d = 1$ we change the transition probability matrix $p_{ij} = p(x_{30} = j | x_{29} = i)$ from the value

$$pT_long(:, :, 29) = \begin{bmatrix} 0.5972 & 0.4028 \\ 0.2999 & 0.7001 \end{bmatrix}$$

to the value

$$pT_long(:, :, 29) = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

With this new value from the transition probability, we can guarantee that regardless the value of x_{29} is, we get $x_{30} = 1$ with probability 1. With this we get that the sequence with highest probability ending in 1 is

>>question12

The sequence with highest probability with $x_d=1$ is:

ans =

Columns 1 through 13

1 1 1 1 1 2 2 2 2 1 2 1 2

Columns 14 through 26

2 2 2 2 1 2 1 1 2 2 2 1 1

Columns 27 through 30

1 1 1 1

Since the most probable chain end with 1, the chain obtained here is exactly the same as the most likely chain, as expected (compare with part 4, exercise 1.1).

5.

To calculate the conditionals $p(x_j | x_{30} = 1)$ we use the product rule to get

$$\mathbb{P}(x_j | x_{30} = 1) = \frac{\mathbb{P}(x_j, x_{30} = 1)}{\mathbb{P}(x_{30} = 1)}.$$

To calculate this number we generate 10000 samples. To calculate the marginals we used the script sampleReject.m shown below

```

1 %Script to do the sample reject to calculate conditionals
2 load viterbiData.mat
3 n=10000;
4
5
6 Y=sampleAncestral(p0,pT_long,n);
7 d=size(Y,2);
8
9
10
11 marginalxd=(sum(Y(:,d)==1)/n);
12
13 aux=Y(:,d)==1;
14 aux=Y(aux,:);
15
16 p1=zeros(d,1);
17 p2=zeros(size(p1));
18 for (k=1:d)
19     temp=aux(:,k)==1;
20     p1temp=sum(temp)/n;
21     temp=aux(:,k)==2;
22     p2temp=sum(temp)/n;
23
24     p1(k)=p1temp/marginalxd;
25     p2(k)=p2temp/marginalxd;
26
27 end
28 display('The number of samples accepted: ')
29 n-size(aux,1)
30
31
32 display('the marginals probabilities are')
33 [p1 p2]

```

Using this script we get the marginals and the acceptance rate to be

```
>>sampleReject
```

The number of samples accepted:

```
ans =
```

```
5618
```

the marginals probabilities are

```
ans =
```

```
0.5961    0.4039
0.4630    0.5370
0.8927    0.1073
```


0.9717	0.0283
0.7713	0.2287
0.2271	0.7729
0.0568	0.9432
0.3366	0.6634
0.3284	0.6716
0.5869	0.4131
0.3676	0.6324
0.7481	0.2519
0.3681	0.6319
0.1376	0.8624
0.3786	0.6214
0.4103	0.5897
0.7241	0.2759
0.7866	0.2134
0.3183	0.6817
0.6593	0.3407
0.5790	0.4210
0.2556	0.7444
0.4147	0.5853
0.2440	0.7560
0.7378	0.2622
0.6120	0.3880
0.5977	0.4023
0.8021	0.1979
0.6392	0.3608
1.0000	0

6.

Now we obtain the conditional marginals $p(x_j|x_{30} = 1)$ using backward sampling. In order to do backward sampling we need to know probabilities of the form $\mathbb{P}(x_i|x_{i+1})$. Since we know all probabilities of the form $\mathbb{P}(x_{i+1}|x_i)$ we can calculate

$$\mathbb{P}(x_i|x_{i+1}) = \frac{\mathbb{P}(x_{i+1}|x_i)\mathbb{P}(x_i)}{\sum_{x_i} \mathbb{P}(x_{i+1}|x_i)\mathbb{P}(x_i)}$$

To do so, we used the script `sampleBackwards.m`. From a 10000 sample we get the following Monte Carlo estimates

```
X=sammplesBackwards(p0,pT_long,10000);
>> [sum(X==1)'/n sum(X==2)'/n]
```

ans =

0.6500	0.3500
0.8670	0.1330
0.9376	0.0624
0.6135	0.3865
0.3802	0.6198
0.0800	0.9200
0.4998	0.5002
0.3060	0.6940
0.5990	0.4010
0.2477	0.7523

0.7289	0.2711
0.4373	0.5627
0.1032	0.8968
0.6182	0.3818
0.8593	0.1407
0.9461	0.0539
0.9083	0.0917
0.3201	0.6799
0.6765	0.3235
0.3067	0.6933
0.2878	0.7122
0.4007	0.5993
0.2568	0.7432
0.4790	0.5210
0.6192	0.3808
0.7010	0.2990
0.7343	0.2657
0.4007	0.5993
0.5825	0.4175
1.0000	0

7.

Finally we calculate the exact conditional marginals $\mathbb{P}(x_j|x_{30}=1)$ by first calculating the functions $M_j(x_j)$ and $V_j(x_j)$ and setting $M_{30}(1) = V_{30} = 1$ along with $M_{30}(2) = V_{30}(2) = 0$. and then we calculate

$$\mathbb{P}(x_j|x_{30} = 1) = \frac{M_j(x_j)V_j(x_j)}{\sum_{x_j} M_j(x_j)V_j(x_j)}.$$

To calculate the marginals we used the script forwardBackwards.m. The results are shown below.

```
>>forwardBackwards(p0,pT_long)
```

```
ans =
```

0.7332	0.2668
0.8495	0.1505
0.9925	0.0075
0.9823	0.0177
0.6826	0.3174
0.0256	0.9744
0.0541	0.9459
0.1854	0.8146
0.4172	0.5828
0.3238	0.6762
0.5978	0.4022
0.6964	0.3036
0.0629	0.9371
0.1944	0.8056
0.7800	0.2200
0.9228	0.0772
0.9607	0.0393
0.6456	0.3544

0.4854	0.5146
0.4718	0.5282
0.3574	0.6426
0.1880	0.8120
0.1859	0.8141
0.2245	0.7755
0.8173	0.1827
0.7944	0.2056
0.8141	0.1859
0.6661	0.3339
0.5625	0.4375
1.0000	0

1.3 1D Linear-Gaussian Markov Chains

Consider a continuous-state Markov chain where the initial distribution is given by

$$x_0 \sim \mathcal{N}(m_0, v_0^2),$$

and the transition distributions for $j > 1$ are given by

$$x_j | x_{j-1} \sim \mathcal{N}(w_j x_{j-1} + m_j, v_j^2).$$

This model could be used to model an object moving through \mathbb{R} .¹ Because of the Gaussian assumptions, this defines a joint Gaussian distribution over the variables while the marginal distributions are also Gaussian. For a generic $j > 1$, derive the form of the marginal distribution of x_j , expressing the marginal parameters μ_j and σ_j recursively in terms of μ_{j-1} and σ_{j-1} .

Hint: You can use Theorem 4.4.1 of Murphy's book.

2 Directed Acyclic Graphical Models

2.1 D-Separation

Consider a directed acyclic graphical (DAG) model with the following graph structure:

Assuming that the conditional independence properties are faithful to the graph, using d-separation [briefly explain why the following are true or false](#):

1. $B \perp F$.
2. $B \perp F \mid A$.
3. $B \perp F \mid C$.
4. $B \perp F \mid E$.
5. $B \perp F \mid I$.

¹In practical applications like object tracking, we typically have that the states x_j are 2- or 3-dimensions so we model an object like a submarine or an airplane moving through space.

6. $B \perp F \mid J$.
7. $B \perp F \mid C, E$.

2.2 Exact Inference

While DAGs can be used as a visual representation of independence assumptions, they can also be used to simplify computations. This question will give you practice using the basic properties which allow efficient computations in graphical models. Consider the DAG model below, for distinguishing between different causes of shortness-of-breath (dyspnoea) and the causes of an abnormal lung x-ray, while modelling potential causes of these diseases too (whether the person is a smoker or had a ‘visit’ to a country with a high degree of tuberculosis).

For this question, let’s assume that we use the following parameterization of the network:

$$\begin{aligned}
 &\text{Visit} \\
 &p(V = 1) = 0.01 \\
 &\text{Smoking} \\
 &p(S = 1) = 0.2 \\
 &\text{Tuberculosis} \\
 &p(T = 1|V = 1) = 0.05 \\
 &p(T = 1|V = 0) = 0.01 \\
 &\text{Lung Cancer} \\
 &p(L = 1|S = 1) = 0.10 \\
 &p(L = 1|S = 0) = 0.01 \\
 &\text{Bronchitis} \\
 &p(B = 1|S = 1) = 0.60 \\
 &p(B = 1|S = 0) = 0.30 \\
 &\text{Abnormal X-Ray} \\
 &p(X = 1|T = 1, L = 1) = 1.00 \\
 &p(X = 1|T = 1, L = 0) = 0.98 \\
 &p(X = 1|T = 0, L = 1) = 0.9 \\
 &p(X = 1|T = 0, L = 0) = 0.05 \\
 \\
 &\text{Dyspnoea} \\
 &p(D = 1|T = 1, L = 1, B = 1) = 0.90 \\
 &p(D = 1|T = 1, L = 1, B = 0) = 0.70 \\
 &p(D = 1|T = 1, L = 0, B = 1) = 0.85 \\
 &p(D = 1|T = 1, L = 0, B = 0) = 0.65 \\
 &p(D = 1|T = 0, L = 1, B = 1) = 0.82 \\
 &p(D = 1|T = 0, L = 1, B = 0) = 0.60 \\
 &p(D = 1|T = 0, L = 0, B = 1) = 0.80 \\
 &p(D = 1|T = 0, L = 0, B = 0) = 0.10
 \end{aligned}$$

Compute the following quantities (hints are given on the right, and these will be easier to do in order and if you use conditional independence properties to simplify the calculations):

0. $p(S = 1)$ (marginal of root node; can read from table)
1. $p(S = 0)$ (negation of marginal of root node; use sum to one constraint)
2. $p(L = 1|S = 1)$ (conditional of child node given parents; can be read from table)
3. $p(L = 1)$ (marginal of child node; marginalize over parent)
4. $p(X = 1|T = 1, L = 1)$ (conditional of child given parents; can be read from table)
5. $p(X = 1|T = 1)$ (conditional of child with missing parent; marginalize over missing parent)
6. $p(X = 1|T = 1, S = 1)$ (conditional of child given parent and grand-parent, marginalize over missing parent)
7. $p(X = 1)$ (marginal of leaf node; marginalize over parents and use independence to simplify)
8. $p(T = 1|X = 1)$ (conditional of parent given child; use Bayes rule)
9. $p(T = 1|L = 1)$ (conditional of parent given co-parent; use independence and then marginal)
10. $p(T = 1|X = 1, L = 1)$ (conditional of parent given child and co-parent; use Bayes rule)

2.3 Inpainting

The function `example_fil.m` loads a variant of the MNIST dataset. It contains all the training images but the test images are missing their bottom half. Running this function fits an independent Bernoulli model to the training set, and then shows the result of applying the density model to “fill in” four random test examples. It performs pretty badly because the independent model can’t condition on the known top-half of the images.

1. Make a variant of the demo where you fit an inhomogeneous Markov chain to each image column. [Hand in your code and an example of using this model to fill in 4 random test images.](#)
2. Make a variant of the demo where you fit a directed acyclic graphical model to the data, using general discrete conditional probabilities and where the parents of pixel (i, j) are the other 8 pixels in the region $(i - 2 : i, j - 2 : j)$. [Hand in your code and an example of using this model to fill in 4 random test images.](#)
3. Consider using more than 8 pixels are parents in the above model, such as the 15 pixels in the region $(i - 3 : i, j - 3 : j)$. If you do this, the code will often place white pixels in the bottom right corner of the image even though no training example has a white pixel there. Why would it do this?
4. Make a variant of the demo where you fit a sigmoid belief network to the data, where the parents of pixel (i, j) are the other pixels in the region $(1 : i, 1 : j)$. [Hand in your code and an example of using this model to fill in 4 random test images.](#)

Hints: For parts 2 and 3, you may find it helpful to make an m by m cell array called `models` where element (i, j) contains the model for pixel (i, j) . For parts 2 and 3 the size of the dataset also mean you will probably need to vectorize your computation. The functions `permute` and `reshape` will help you, making a sparse version of X with `sparse` can also speed up many operations. For part 2, you can use `binaryTabular.m` to fit the discrete conditional distribution and sample from it (a reasonable value of α is 1). For part 3, you can use `logisticL2.m` to fit logistic regression models and sample from them (a reasonable value of λ is 1). Note that `logisticL2.m` uses a $\{-1, 1\}$ encoding of y while `binaryTabular.m` uses a $\{0, 1\}$ encoding (both support sparse X).

Solutions

1.

First we fitted an inhomogenous Markov chain, where given a pixel x_j we calculated the transition probability from x_{j-1} as

$$\mathbb{P}(x_j|x_{j-1}) = \frac{\text{Number of times we have a transition from } x_{j-1} \text{ to } x_j}{\text{Number of times we started at } x_j}.$$

We trained our model using the script `demoInhomogeneousMC.m`. We used the model on 4 random images from the test set. The results are shown below




Figure 1: Four samples from an inhomogeneous Markov Chain model

2 and 3.

In this part of the exercise we trained two DAGs where each pixel had 8 parents and 15 partents. The parents for the pixel (i, j) where the pixels in the region $((i - 2) : i, (j - 2) : j)$ and $((i - 3) : i, (j - 3) : j)$ repectively. The scripts used to create the model are `demoDAG8.m` and `demoDAG15.m`. The results from this different models are shown below.




Figure 2: Four sample images from a DAG model with 8 parents.

The results for the DAG with 15 parents are




Figure 3: Four sample images from a DAG model with 15 parents.

It is interesting to see that often the pixel located at $(28, 28)$ has value 1 despite all training images have the value of 0 at that location. To understand why this is case consider the parent located at $(25, 25)$. There are 16 training images with 1 in that entry also pixels $(25, 26)$ has five images with that characteristic and $(26, 25)$ five images. Hence there are different diffent combinations of vectors $x_{pa((28,28))} \in \{0, 1\}^{15}$ such that

$$\mathbb{P}(x_{(28,28)}=1|x_{pa((28,28))}) \neq 0.$$

4. Finally we trained a sigmoid belief network where the parents of the pixel located at the point (i, j) are the pixels in the region $(1 : i, 1 : j)$.