

CPSC 540: Machine Learning

More DAGs, Undirected Graphical Models

Mark Schmidt

University of British Columbia

Winter 2017

Admin

- **Assignment 3:**
 - 2 late days to hand in today.
- **Assignment 4:**
 - Due March 20.
- For **graduate students planning to graduate in May:**
 - Send me a private message on Piazza ASAP.

Last Time: Directed Acyclic Graphical (DAG) Models

- DAG models use a factorization of the joint distribution,

$$p(x_1, x_2, \dots, x_d) = \prod_{j=1}^d p(x_j | x_{\text{pa}(j)}),$$

where $\text{pa}(j)$ are the **parents** of node j .

- This assumes a **Markov property**,

$$p(x_j | x_{1:j-1}) = p(x_j | x_{\text{pa}(j)}),$$

which **generalizes the Markov property in Markov chains**,

$$p(x_j | x_{1:j-1}) = p(x_j | x_{j-1}).$$


Last Time: Directed Acyclic Graphical (DAG) Models

- DAG models use a factorization of the joint distribution,

$$p(x_1, x_2, \dots, x_d) = \prod_{j=1}^d p(x_j | x_{\text{pa}(j)}),$$

where $\text{pa}(j)$ are the **parents** of node j .

- We visualize the assumptions made by the model as a graph:



- Structure determines **conditional independences** and **computational tractability**.

Outline

- 1 D-Separation and Plate Notation
- 2 Learning and Inference in DAGs
- 3 Undirected Graphical Models

D-Separation

- We say that A and B are **d-separated** (conditionally independent) if *all paths* P from A to B are “blocked” because *at least one* of the following holds:


- ① P includes a “chain” with an observed middle node (e.g., Markov chain):



- ② P includes a “fork” with an observed parent node (e.g., mixture model):




- ③ P includes a “v-structure” or “collider” (e.g., factor analysis):



where “child” and all its descendants are unobserved.

Alarm Example



- Earthquake $\not\perp$ Call.
- Earthquake \perp Call | Alarm.
- Alarm $\not\perp$ Stuff Missing.
- Alarm \perp Stuff Missing | Burglary.

Alarm Example



- Earthquake \perp Burglary.
- Earthquake $\not\perp$ Burglary | Alarm.
 - **Explaining away**: Knowing Earthquake would make Burglary is less likely.
- Call $\not\perp$ Stuff Missing.
- Earthquake \perp Stuff Missing.
- Earthquake $\not\perp$ Stuff Missing | Call.

Discussion of D-Separation

- D-separation lets you say if **conditional independence is implied** by assumptions:

$$(A \text{ and } B \text{ are d-separated given } E) \Rightarrow A \perp B \mid E.$$

- However, there **might be extra conditional independences** in the distribution:

- These would depend on specific choices of the $p(x_j | x_{\text{pa}(j)})$.
- Or some *orderings* may reveal extra independences....

- Instead of restricting to $\{1, 2, \dots, j-1\}$, consider **general parent choices**.


- x_2 could be a parent of x_1 .

- As long the **graph is acyclic**, there exists a valid ordering.

(all DAGs have a “topological order” of variables where parents are before children)

Non-Uniqueness of Graph and Equivalent Graphs

- Note that some graphs imply **same conditional independences**:
 - Equivalent** graphs: same v-structures and other (undirected) edges are the same.
 - Examples of 3 *equivalent* graphs (left) and 3 non-equivalent graphs (right):




Discussion of D-Separation

- So the graph is not necessarily unique and is not the whole story.
- But, we can do a lot with d-separation:
 - Implies every independence/conditional-independence we've used in 340/540.
- Here we start blurring distinction between data/parameters/hyper-parameters...

IID Assumption as a DAG


- On Day 2, our first independence assumption was the **IID assumption**:




- Training/test examples come independently from data-generating process D .
- If we knew D , then there would be no need to learn.
- But D is **unobserved**, so knowing about some x^i tells us about the others.
- We'll use this understanding later to **relax the IID assumption**.

Plate Notation

- Graphical representation of the IID assumption:



- We can concisely represent repeated parts of graphs using plate notation:




Tilde Notation as a DAG


- When we write

$$y^i \sim \mathcal{N}(w^T x^i, 1),$$


we can interpret it as the DAG model:



- If the x^i are IID then we can represent supervised learning as



or




- From d -separation on this graph we have $p(y|X, w) = \prod_{i=1}^n p(y^i|x^i, w)$.

Tilde Notation as a DAG


- When we do MAP estimation under the assumptions

$$y^i \sim \mathcal{N}(w^T x^i, 1), \quad w_j \sim \mathcal{N}(0, 1/\lambda),$$

we can interpret it as the DAG model:




- Or introducing a second plate using:




Other Models in DAG/Plate Notation

- For naive Bayes or Gaussian discriminant analysis with diagonal Σ_c we have

$$y^i \sim \text{Cat}(\theta), \quad x^i | y^i = c \sim D(\theta_c).$$



- Or in plate notation as



Other Models in DAG/Plate Notation


- In a full Gaussian model for a single x we have

$$x^i \sim \mathcal{N}(\mu, \Sigma).$$



- For mixture of Gaussians we have

$$z^i \sim \text{Cat}(\theta), \quad x^i | z^i = c \sim \mathcal{N}(\mu_c, \Sigma_c).$$



Outline

- 1 D-Separation and Plate Notation
- 2 Learning and Inference in DAGs
- 3 Undirected Graphical Models

Parameter Learning in General DAG Models

- The log-likelihood in DAG models is **separable** in the conditionals,

$$\begin{aligned}\log p(x|\Theta) &= \log \prod_{j=1}^d p(x_j|x_{\text{pa}(j)}, \Theta_j) \\ &= \sum_{j=1}^d \log p(x_j|x_{\text{pa}(j)}, \Theta_j)\end{aligned}$$

- If each $p(x_j|x_{\text{pa}(j)})$ has its own parameters Θ_j , we can **fit them independently**.
 - We've done this before: naive Bayes, Gaussian discriminant analysis, etc.
- Sometimes you want to have **tied parameters** ($\Theta_j = \Theta_{j'}$)
 - Homogeneous Markov chains, Gaussian discriminant analysis with shared covariance.
 - Still easy, but need to fit $p(x_j|x_{\text{pa}(j)}, \Theta_j)$ and $p(x_{j'}|x_{\text{pa}(j')}, \Theta_j)$ together.

Tabular Parameterization in DAG Models

- To specify distribution, we need to decide on the form of $p(x_j|x_{\text{pa}(j)}, \Theta_j)$.
- For discrete data a default choice is the **tabular parameterization**:

$$p(x_j|x_{\text{pa}(j)}, \Theta_j) = \theta_{x_j, x_{\text{pa}(j)}},$$


as we did for Markov chains (but now with multiple parents).

- **Intuitive**: just need conditional probabilities of children given parents like

$$p(\text{"wet grass"} = 1 \mid \text{"sprinkler"} = 1, \text{"rain"} = 0),$$

and MLE is just counting.

Tabular Parameterization Example




https://en.wikipedia.org/wiki/Bayesian_network

Some quantities can be directly read from the tables:

$$p(R = 1) = 0.2.$$

$$p(G = 1|S = 0, R = 1) = 0.8.$$

Tabular Parameterization Example




https://en.wikipedia.org/wiki/Bayesian_network

Can calculate any probabilities using marginalization/product-rule/Bayes-rule.

$$\begin{aligned}
 p(G = 1|R = 1) &= p(G = 1, S = 0|R = 1) + p(G = 1, S = 1|R = 1) & \left(p(a|c) = \sum_b p(a, b|c) \right) \\
 &= p(G = 1|S = 0, R = 1)p(S = 0|R = 1) + p(G = 1|S = 1, R = 1)p(S = 1|R = 1) \\
 &= 0.8(0.99) + 0.99(0.01) = 0.81.
 \end{aligned}$$

Tabular Parameterization Example

Some companies sell software to help companies reason using tabular DAGs:



Fitting DAGs using Supervised Learning


- But tabular parameterization requires **too many parameters**:
 - With binary states and k parents, need 2^{k+1} **parameters**.
- One solution is letting users specify a “parsimonious” parameterization:
 - Typically have a linear number of parameters.
 - For example, the “noisy-or” model: $p(x_j | x_{\text{pa}(j)}) = 1 - \prod_{k \in \text{pa}(j)} q_j$.
- But if we have data, we can use **supervised learning**.
 - Write fitting $p(x_j | x_{\text{pa}(j)})$ as our usual $p(y|x)$.
 - We’re **predicting one column of X given the values of other columns**.

Fitting DAGs using Supervised Learning

- Fitting DAGs using supervised learning:
 - For $j = 1 : d$:
 - 1 Set $\tilde{y}^i = x_j^i$ and $\tilde{x}^i = x_{\text{pa}(j)}^i$.
 - 2 Solve a supervised learning problem using $\{\tilde{X}, \tilde{y}\}$.
 - Use the d regression/classification models as the density estimator.
- We can use our usual tricks:
 - Linear models, non-linear bases, regularization, kernel trick, random forests, etc.
 - With least squares it's called a Gaussian belief network.
 - With logistic regression it's called a sigmoid belief networks.
 - Don't need Markov assumptions to tractably fit these models.

MNIST Digits with Tabular DAG Model

- Recall our latest MNIST model using a **tabular DAG**:




- This model is pretty bad because you only see 8 parents.

MNIST Digits with Sigmoid Belief Network

- Samples from **sigmoid belief network**:

(DAG with logistic regression for each variable)



where we use all previous pixels as parents (from 0 to 783 parents).

- Models **long-range dependencies** but has a **linear assumption**.

Sampling in DAGs

- We can use **ancestral sampling** to generate samples from a DAG:
 - 1 Sample x_1 from $p(x_1)$.
 - 2 If x_1 is a parent of x_2 , sample x_2 from $p(x_2|x_1)$.
 - Otherwise, sample x_2 from $p(x_2)$.
 - 3 Go through the subsequent j in order sampling x_j from $p(x_j|x_{\text{pa}(j)})$.
- We can use these samples within **Monte Carlo** methods.
- How do **sample from a multivariate Gaussian**?
 - Write it as a Gaussian belief network, apply ancestral sampling.


Inference in Forest DAGs

- If we try to generalize the **CK equations** to DAGs we obtain

$$p(x_j = s) = \sum_{x_{\text{pa}(j)}} p(x_j = s, x_{\text{pa}(j)}) = \sum_{x_{\text{pa}(j)}} \underbrace{p(x_j = s | x_{\text{pa}(j)})}_{\text{given}} p(x_{\text{pa}(j)}).$$

which **works if each node has at most one parent**.

- Such graphs are called **trees** (connected), or **forests** (disconnected).
 - Also called “singly-connected”.




- **Forests allow efficient message-passing** methods as in Markov chains.
 - E.g., decoding and computing univariate marginals/conditionals in $O(dk^2)$.
 - Message passing applied to tree-structured graphs is called **belief propagation**.

Inference in General DAGs

- If we try to generalize the **CK equations** to DAGs we obtain

$$p(x_j = s) = \sum_{x_{\text{pa}(j)}} p(x_j = s, x_{\text{pa}(j)}) = \sum_{x_{\text{pa}(j)}} \underbrace{p(x_j = s | x_{\text{pa}(j)})}_{\text{given}} p(x_{\text{pa}(j)}).$$

- What goes wrong if nodes have multiple parents?
 - The expression $p(x_{\text{pa}(j)})$ is a **joint distribution** and is not given recursively.
- Consider the non-tree graph:



Inference in General DAGs

- We can compute $p(x_4)$ in this non-tree using:

$$\begin{aligned} p(x_4) &= \sum_{x_3} \sum_{x_2} \sum_{x_1} p(x_1, x_2, x_3, x_4) \\ &= \sum_{x_3} \sum_{x_2} \sum_{x_1} p(x_4|x_2, x_3)p(x_3|x_1)p(x_2|x_1)p(x_1) \\ &= \sum_{x_3} \sum_{x_2} p(x_4|x_2, x_3) \underbrace{\sum_{x_1} p(x_3|x_1)p(x_2|x_1)p(x_1)}_{M_{23}(x_2, x_3)} \end{aligned}$$

- Dependencies between $\{x_1, x_2, x_3\}$ mean our **message depends on two variables**.


$$\begin{aligned} p(x_4) &= \sum_{x_3} \sum_{x_2} p(x_4|x_2, x_3)M_{23}(x_2, x_3) \\ &= \sum_{x_3} M_{34}(x_3, x_4), \end{aligned}$$

Inference in General DAGs


- With 2-variable messages, our **cost increases** to $O(dk^3)$.
- If we add the edge $x_1 -> x_4$, then the cost is $O(dk^4)$.
(the same cost as enumerating all possible assignments)
- Unfortunately, cost is **not as simple as counting number of parents**.
 - Even if each node has 2 parents, we may need huge messages.
 - Decoding is NP-hard and marginals are #P-hard in general.
 - We'll see later that maximum message is given by **treewidth** of a particular graph.
- In general, we'll need **approximate inference** methods to use general DAGs.

Conditional Sampling in DAGs

- What about **conditional sampling** in DAGs?
 - Could be easy or hard depending on what we condition on.
- For example, still **easy if we condition on the first** variables in the order:
 - Just fix these and run ancestral sampling.



- **Hard to condition on the last** variables in the order:
 - Conditioning on descendent makes ancestors dependent.



Outline

- 1 D-Separation and Plate Notation
- 2 Learning and Inference in DAGs
- 3 Undirected Graphical Models**

Directed vs. Undirected Models

- In some applications we have a **natural ordering** of the x_j .
 - In the “rain” data, the past affects the future.
- In some applications we **don't have a natural order**.
 - E.g., pixels in an image.
- In these settings we often use **undirected graphical models**.
 - Also known as **Markov random fields** and originally from statistical physics.

Undirected Graphical Models

- Undirected graphical models (UGMs) assume $p(x)$ factorizes over subsets c ,

$$p(x_1, x_2, \dots, x_d) \propto \prod_{c \in \mathcal{C}} \phi_c(x_c),$$

from among a set of subsets of \mathcal{C} .


- The ϕ_c are called **potential functions**: can be **any non-negative function**.
 - **Ordering doesn't matter**: more natural for things like pixels of an image.
 - Theoretically, only need ϕ_c for maximal subsets in \mathcal{C} .
- Important special case is **pairwise** undirected graphical model:

$$p(x) \propto \left(\prod_{j=1}^d \phi_j(x_j) \right) \left(\prod_{(i,j) \in E} \phi_{ij}(x_i, x_j) \right),$$

where E are a set of undirected edges.

Undirected Graphical Models

- Pairwise UGMs are a classic way to model dependencies in images:



- Can model dependency between neighbouring pixels, without imposing ordering.

From Probability Factorization to Graphs

- For a pairwise UGM,

$$p(x) \propto \left(\prod_{j=1}^d \phi_j(x_j) \right) \left(\prod_{(i,j) \in E} \phi_{ij}(x_i, x_j) \right),$$

we visualize independence assumptions as an undirected graph:

- We have edge i to j if $(i, j) \in E$.


- For general UGMs,

$$p(x_1, x_2, \dots, x_d) \propto \prod_{c \in \mathcal{C}} \phi_c(x_c),$$

we have the edge (i, j) if i and j are together in at least one c .

Conditional Independence in Undirected Graphical Models

- It's easy to check **conditional independence** in UGMs:
 - $A \perp B | C$ if C **blocks all paths** from any A to any B .
- Example:



- $A \not\perp C$.
- $A \not\perp C | B$.
- $A \perp C | B, E$.
- $A, B \not\perp F | C$
- $A, B \perp F | C, E$.

Multivariate Gaussian and Pairwise Graphical Models

- **Multivariate Gaussian** is a special case of a **pairwise UGM**.
- Edges of the graph are (i, j) values where $\Sigma_{ij}^{-1} \neq 0$.
- Unconditional independence of (i, j) corresponds to having $\Sigma_{ij} = 0$.
 - Can be seen from block Gaussian formula.
 - Corresponds to **reachability in the graph**.
- We use the term **Gaussian graphical model** (GGM) in this context.
 - Or **Gaussian Markov random field**.

Digression: Gaussian Graphical Models

- Multivariate Gaussian can be written as

$$p(x) \propto \exp\left(-\frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu)\right) \propto \exp\left(-\frac{1}{2}x^T \Sigma^{-1}x + x^T \underbrace{\Sigma^{-1}\mu}_v\right),$$

and from here we can see that it's a **pairwise UGM**:

$$\begin{aligned} p(x) &\propto \exp\left(-\frac{1}{2} \sum_{i=1}^d \sum_{j=1}^d x_i x_j \Sigma_{ij}^{-1} + \sum_{i=1}^d x_i v_i\right) \\ &= \left(\prod_{i=1}^d \prod_{j=1}^d \underbrace{\exp\left(-\frac{1}{2} x_i x_j \Sigma_{ij}^{-1}\right)}_{\phi_{ij}(x_i, x_j)} \right) \left(\prod_{i=1}^d \underbrace{\exp(x_i v_i)}_{\phi_i(x_i)} \right) \end{aligned}$$

Independence in GGMs

- So Gaussians are pairwise UGMs with $\phi_{ij}(x_i, x_j) = \exp\left(-\frac{1}{2}x_i x_j \Theta_{ij}\right)$,
 - Where Θ_{ij} is element (i, j) of Σ^{-1} .
- Connection between **precision matrix $\Theta = \Sigma^{-1}$** and **conditional independence**:
 - Setting **$\Theta_{ij} = 0$** is equivalent to removing $\phi_{ij}(x_i, x_j)$ from the UGM.

$$\Theta_{ij} \neq 0 \Rightarrow x_i \not\perp x_j | x_{-ij}.$$

- Gaussian **conditional independencies corresponds to sparsity in precision matrix**.
 - Diagonal Θ gives disconnected graph: all variables are independent.
 - Full Θ gives fully-connected graph: there are no independences.

Independence in GGMs

- Consider Gaussian with **tri-diagonal precision** Θ :

$$\Sigma^{-1} = \begin{bmatrix} 32.0897 & 13.1740 & 0 & 0 & 0 \\ 13.1740 & 18.3444 & -5.2602 & 0 & 0 \\ 0 & -5.2602 & 7.7173 & 2.1597 & 0 \\ 0 & 0 & 2.1597 & 20.1232 & 1.1670 \\ 0 & 0 & 0 & 1.1670 & 3.8644 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} 0.0494 & -0.0444 & -0.0312 & 0.0034 & -0.0010 \\ -0.0444 & 0.1083 & 0.0761 & -0.0083 & 0.0025 \\ -0.0312 & 0.0761 & 0.1872 & -0.0204 & 0.0062 \\ 0.0034 & -0.0083 & -0.0204 & 0.0528 & -0.0159 \\ -0.0010 & 0.0025 & 0.0062 & -0.0159 & 0.2636 \end{bmatrix}$$

- $\Sigma_{ij} \neq 0$ so all variables are dependent: $x_1 \not\perp x_2$, $x_1 \not\perp x_5$, and so on.
- But **conditional independence is described by a Markov chain**:

$$p(x_1 | x_2, x_3, x_4, x_5) = p(x_1 | x_2).$$

Graphical Lasso

- Conditional independence in GGMs is described by sparsity in Θ .
- Recall fitting multivariate Gaussian with L1-regularization,

$$\operatorname{argmin}_{\Theta \succ 0} \operatorname{Tr}(S\Theta) - \log |\Theta| + \lambda \|\Theta\|_1,$$

which is called the **graphical Lasso** because it **encourages a sparse graph**.

- Special case of graph **structure learning**.
- Consider instead fitting DAG model with Gaussian probabilities:
 - **DAG structure corresponds to sparsity in Cholesky** of covariance.

Tractability of UGMs

- In UGMs we assume that

$$p(x) = \frac{1}{Z} \prod_{c \in \mathcal{C}} \phi_c(x_c),$$

where Z is the constant such that

$$\sum_{x_1} \sum_{x_2} \cdots \sum_{x_d} p(x) = 1 \text{ (discrete),} \quad \int_{x_1} \int_{x_2} \cdots \int_{x_d} p(x) dx_d dx_{d-1} \cdots dx_1 = 1 \text{ (cont).}$$

- So Z is

$$Z = \sum_x \prod_{c \in \mathcal{C}} \phi_c(x_c) \text{ (discrete),} \quad \int_x \prod_{c \in \mathcal{C}} \phi_c(x_c) dx \text{ (cont)}$$

- Whether you can compute Z depends on the choice of ϕ_c :
 - Gaussian case: $O(d^3)$ in general, but $O(d)$ for forests (no loops).
 - Discrete case: #P-hard in general, but $O(dk^2)$ for forests (no loops).
 - Continuous non-Gaussian: usually requires numerical integration.

Summary

- **Plate Notation** lets compactly draw graphs with repeated patterns.
 - There are fancier versions of plate notation called “probabilistic programming”.
- **Parameter learning in DAGs:**
 - Can fit each $p(x_j | x_{\text{pa}(j)})$ independently.
 - Tabular parameterization, or treat as supervised learning.
- **Inference in DAGs:**
 - Ancestral sampling and Monte Carlo methods work as faster.
 - Message-passing message sizes depend on graph structure.
- **Undirected graphical models** factorize probability into non-negative potentials.
 - Simple conditional independence properties.
 - Include Gaussians as special case.
- Next time: our first visit to the wild world of approximate inference.